

## 6\_0\_2\_Model\_Evaluation\_and\_Error\_Analysis—Multi

April 13, 2025

### 0.1 Packages, Library Imports, File Mounts, & Data Imports \*\* Run All \*\*

```
[ ]: !pip install -q transformers
!pip install -q torchinfo
!pip install -q datasets
!pip install -q evaluate
!pip install -q nltk
!pip install -q contractions
!pip install -q hf_xet
!pip install -q sentencepiece
# !pip install -q import openpyxl
```

```
[ ]: !sudo apt-get update
! sudo apt-get install tree
```

```
Hit:1 https://cloud.r-project.org/bin/linux/ubuntu jammy-cran40/ InRelease
Hit:2 http://security.ubuntu.com/ubuntu jammy-security InRelease
Hit:3 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64
InRelease
Hit:4 https://r2u.stat.illinois.edu/ubuntu jammy InRelease
Hit:5 http://archive.ubuntu.com/ubuntu jammy InRelease
Hit:6 http://archive.ubuntu.com/ubuntu jammy-updates InRelease
Hit:7 https://ppa.launchpadcontent.net/deadsnakes/ppa/ubuntu jammy InRelease
Hit:8 http://archive.ubuntu.com/ubuntu jammy-backports InRelease
Hit:9 https://ppa.launchpadcontent.net/graphics-drivers/ppa/ubuntu jammy
InRelease
Hit:10 https://ppa.launchpadcontent.net/ubuntugis/ppa/ubuntu jammy InRelease
Reading package lists... Done
W: Skipping acquire of configured file 'main/source/Sources' as repository
'https://r2u.stat.illinois.edu/ubuntu jammy InRelease' does not seem to provide
it (sources.list entry misspelt?)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
tree is already the newest version (2.0.2-1).
0 upgraded, 0 newly installed, 0 to remove and 31 not upgraded.
```

```
[ ]: #@title Imports
import nltk
from nltk.tokenize import RegexpTokenizer
import sentencepiece
import contractions
import spacy

import evaluate
from datasets import load_dataset, Dataset, DatasetDict

import torch
import torch.nn as nn
from torchinfo import summary

import transformers
from transformers import AutoTokenizer, AutoModel, \
    ↳AutoModelForSequenceClassification, TrainingArguments, Trainer, BertConfig, \
    ↳BertForSequenceClassification

import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import sklearn
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import classification_report, \
    ↳precision_recall_fscore_support, accuracy_score

import json
import datetime
import zoneinfo
from datetime import datetime

import math
from mpl_toolkits.mplot3d import Axes3D
from sklearn.decomposition import PCA
from scipy.stats import entropy
from sklearn.metrics import confusion_matrix

[ ]: # @title Mount Google Drive

[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[ ]: dir_root = '/content/drive/MyDrive/266-final/'  
# dir_data = '/content/drive/MyDrive/266-final/data/'  
# dir_data = '/content/drive/MyDrive/266-final/data/se21-t1-comp-lex-master/'  
dir_data = '/content/drive/MyDrive/266-final/data/266-comp-lex-master'  
dir_models = '/content/drive/MyDrive/266-final/models/'  
dir_results = '/content/drive/MyDrive/266-final/results/'  
log_filename = "experiment_runs.txt"  
log_filepath = os.path.join(dir_results, log_filename)
```

```
[ ]: wandbai_api_key = ""
```

```
[ ]: !tree /content/drive/MyDrive/266-final/data/266-comp-lex-master/
```

```
/content/drive/MyDrive/266-final/data/266-comp-lex-master/  
  fe-test-labels  
    test_multi_df.csv  
    test_single_df.csv  
  fe-train  
    train_multi_df.csv  
    train_single_df.csv  
  fe-trial-val  
    trial_val_multi_df.csv  
    trial_val_single_df.csv  
  test-labels  
    lcp_multi_test.tsv  
    lcp_single_test.tsv  
  train  
    lcp_multi_train.tsv  
    lcp_single_train.tsv  
  trial  
    lcp_multi_trial.tsv  
    lcp_single_trial.tsv
```

6 directories, 12 files

```
[ ]: !ls -R /content/drive/MyDrive/266-final/data/266-comp-lex-master/
```

```
/content/drive/MyDrive/266-final/data/266-comp-lex-master/:  
fe-test-labels  fe-train  fe-trial-val  test-labels  train  trial  
  
/content/drive/MyDrive/266-final/data/266-comp-lex-master/fe-test-labels:  
test_multi_df.csv  test_single_df.csv  
  
/content/drive/MyDrive/266-final/data/266-comp-lex-master/fe-train:  
train_multi_df.csv  train_single_df.csv
```

```
/content/drive/MyDrive/266-final/data/266-comp-lex-master/fe-trial-val:  
trial_val_multi_df.csv trial_val_single_df.csv
```

```
/content/drive/MyDrive/266-final/data/266-comp-lex-master/test-labels:  
lcp_multi_test.tsv lcp_single_test.tsv
```

```
/content/drive/MyDrive/266-final/data/266-comp-lex-master/train:  
lcp_multi_train.tsv lcp_single_train.tsv
```

```
/content/drive/MyDrive/266-final/data/266-comp-lex-master/trial:  
lcp_multi_trial.tsv lcp_single_trial.tsv
```

```
[ ]: !tree /content/drive/MyDrive/266-final/data/266-comp-lex-master/
```

```
/content/drive/MyDrive/266-final/data/266-comp-lex-master/  
  fe-test-labels  
    test_multi_df.csv  
    test_single_df.csv  
  fe-train  
    train_multi_df.csv  
    train_single_df.csv  
  fe-trial-val  
    trial_val_multi_df.csv  
    trial_val_single_df.csv  
  test-labels  
    lcp_multi_test.tsv  
    lcp_single_test.tsv  
  train  
    lcp_multi_train.tsv  
    lcp_single_train.tsv  
  trial  
    lcp_multi_trial.tsv  
    lcp_single_trial.tsv
```

6 directories, 12 files

```
[ ]: #@title Import Data
```

```
[ ]: df_names = [  
    "train_single_df",  
    "train_multi_df",  
    "trial_val_single_df",  
    "trial_val_multi_df",  
    "test_single_df",  
    "test_multi_df"  
]  
  
loaded_dataframes = {}
```

```

for df_name in df_names:
    if "train" in df_name:
        subdir = "fe-train"
    elif "trial_val" in df_name:
        subdir = "fe-trial-val"
    elif "test" in df_name:
        subdir = "fe-test-labels"
    else:
        subdir = None

    if subdir:
        read_path = os.path.join(dir_data, subdir, f"{df_name}.csv")
        loaded_df = pd.read_csv(read_path)
        loaded_dataframes[df_name] = loaded_df
        print(f"Loaded {df_name} from {read_path}")

# for df_name, df in loaded_dataframes.items():
#     print(f"\n>>> {df_name} shape: {df.shape}")
#     if 'binary_complexity' in df.columns:
#         print(df['binary_complexity'].value_counts())
#         print(df.info())
#         print(df.head())

for df_name, df in loaded_dataframes.items():
    globals()[df_name] = df
    print(f"{df_name} loaded into global namespace.")

```

```

Loaded train_single_df from /content/drive/MyDrive/266-final/data/266-comp-lex-
master/fe-train/train_single_df.csv
Loaded train_multi_df from /content/drive/MyDrive/266-final/data/266-comp-lex-
master/fe-train/train_multi_df.csv
Loaded trial_val_single_df from /content/drive/MyDrive/266-final/data/266-comp-
lex-master/fe-trial-val/trial_val_single_df.csv
Loaded trial_val_multi_df from /content/drive/MyDrive/266-final/data/266-comp-
lex-master/fe-trial-val/trial_val_multi_df.csv
Loaded test_single_df from /content/drive/MyDrive/266-final/data/266-comp-lex-
master/fe-test-labels/test_single_df.csv
Loaded test_multi_df from /content/drive/MyDrive/266-final/data/266-comp-lex-
master/fe-test-labels/test_multi_df.csv
train_single_df loaded into global namespace.
train_multi_df loaded into global namespace.
trial_val_single_df loaded into global namespace.
trial_val_multi_df loaded into global namespace.
test_single_df loaded into global namespace.
test_multi_df loaded into global namespace.

```

- Functional tests pass, we can proceed with Baseline Modeling

## 0.2 Experiments

### 0.2.1 Helper Functions **\*\* Run \*\***

```
[ ]: MODEL_LINEAGE = {}

def get_model_and_tokenizer(
    remote_model_name: str = None,
    local_model_path: str = None,
    config=None
):
    """
    Loads the model & tokenizer for classification.
    If 'local_model_path' is specified, load from that path.
    Otherwise, fall back to 'remote_model_name'.

    Optional: 'config' can be a custom BertConfig/AutoConfig object
    to override certain configuration parameters.

    Records complete traceable lineage in the global MODEL_LINEAGE.
    """
    global MODEL_LINEAGE

    if local_model_path:
        print(f>Loading from local path: {local_model_path}")
        tokenizer = AutoTokenizer.from_pretrained(local_model_path)

        # If a config object is provided, we pass it to from_pretrained.
        # Otherwise, it just uses the config that is part of local_model_path.
        if config is not None:
            model = AutoModelForSequenceClassification.from_pretrained(
                local_model_path,
                config=config
            )
        else:
            model = AutoModelForSequenceClassification.
↳from_pretrained(local_model_path)

        MODEL_LINEAGE = {
            "type": "offline_checkpoint",
            "path": local_model_path,
            "timestamp": datetime.now().strftime("%Y-%m-%d %H:%M:%S")
        }
    elif remote_model_name:
        print(f>Loading from Hugging Face model: {remote_model_name}")
        tokenizer = AutoTokenizer.from_pretrained(remote_model_name)

        if config is not None:
```

```

        model = AutoModelForSequenceClassification.from_pretrained(
            remote_model_name,
            config=config
        )
    else:
        model = AutoModelForSequenceClassification.
↳from_pretrained(remote_model_name)

    MODEL_LINEAGE = {
        "type": "huggingface_hub",
        "path": remote_model_name,
        "timestamp": datetime.now().strftime("%Y-%m-%d %H:%M:%S")
    }
    else:
        raise ValueError("You must provide either a remote_model_name or a
↳local_model_path!")

    return model, tokenizer

```

```

[ ]: def freeze_unfreeze_layers(model, layers_to_unfreeze=None):
    """
    Toggles requires_grad = False for all parameters
    except for those whose names contain any string in layers_to_unfreeze.
    By default, always unfreeze classifier/heads.
    """
    if layers_to_unfreeze is None:
        layers_to_unfreeze = ["classifier.", "pooler."]

    for name, param in model.named_parameters():
        if any(substring in name for substring in layers_to_unfreeze):
            param.requires_grad = True
        else:
            param.requires_grad = False

```

```

[ ]: def encode_examples(examples, tokenizer, text_col, max_length=256):
    """
    Tokenizes a batch of texts from 'examples[text_col]' using the given
↳tokenizer.
    Returns a dict with 'input_ids', 'attention_mask', etc.
    """
    texts = examples[text_col]
    encoded = tokenizer(
        texts,
        truncation=True,
        padding='max_length',
        max_length=max_length
    )

```

```
return encoded
```

```
[ ]: def prepare_dataset(df, tokenizer, text_col, label_col, max_length=256):  
    """  
    Converts a Pandas DataFrame to a Hugging Face Dataset,  
    then applies 'encode_examples' to tokenize.  
    """  
    dataset = Dataset.from_pandas(df)  
  
    dataset = dataset.map(  
        lambda batch: encode_examples(batch, tokenizer, text_col, max_length),  
        batched=True  
    )  
  
    dataset = dataset.rename_column(label_col, "labels")  
    dataset.set_format(type='torch',  
                        columns=['input_ids', 'attention_mask', 'labels'])  
    return dataset  
  
[ ]: def compute_metrics(eval_pred):  
    """  
    Computes classification metrics, including accuracy, precision, recall, and  
    ↪F1.  
    """  
    logits, labels = eval_pred  
    preds = np.argmax(logits, axis=1)  
  
    metric_accuracy = evaluate.load("accuracy")  
    metric_precision = evaluate.load("precision")  
    metric_recall = evaluate.load("recall")  
    metric_f1 = evaluate.load("f1")  
  
    accuracy_result = metric_accuracy.compute(predictions=preds, ↪  
    ↪references=labels)  
    precision_result = metric_precision.compute(predictions=preds, ↪  
    ↪references=labels, average="binary")  
    recall_result = metric_recall.compute(predictions=preds, ↪  
    ↪references=labels, average="binary")  
    f1_result = metric_f1.compute(predictions=preds, references=labels, ↪  
    ↪average="binary")  
  
    return {  
        "accuracy" : accuracy_result["accuracy"],  
        "precision": precision_result["precision"],  
        "recall" : recall_result["recall"],  
        "f1" : f1_result["f1"]  
    }
```



```
[ ]: def gather_config_details(model):
    """
    Enumerates every attribute in model.config
    """
    config_items = {}
    for attr_name, attr_value in vars(model.config).items():
        config_items[attr_name] = attr_value
    return config_items

def gather_model_details(model):
    """
    Extracts total layers, total params, trainable params, and activation_
    ↪function
    from a Transformers model. Adjust logic as needed for different_
    ↪architectures.
    """
    details = {}

    try:
        total_params = model.num_parameters()
        trainable_params = model.num_parameters(only_trainable=True)
    except AttributeError:
        all_params = list(model.parameters())
        total_params = sum(p.numel() for p in all_params)
        trainable_params = sum(p.numel() for p in all_params if p.requires_grad)

    details["model_total_params"] = total_params
    details["model_trainable_params"] = trainable_params

    if hasattr(model, "bert") and hasattr(model.bert, "pooler"):
        act_obj = getattr(model.bert.pooler, "activation", None)
        details["pooler_activation_function"] = act_obj.__class__.__name__ if_
    ↪act_obj else "N/A"
    else:
        details["pooler_activation_function"] = "N/A"

    details["config_attributes"] = gather_config_details(model)
    return details

def gather_all_run_metrics(trainer, train_dataset=None, val_dataset=None,
    ↪test_dataset=None):
    """
    Gathers final training metrics, final validation metrics, final test_
    ↪metrics.
    Instead of only parsing the final train_loss from the log, we also do a full
    trainer.evaluate(train_dataset) to get the same set of metrics that val/
    ↪test have.

```

```

"""
results = {}

if train_dataset is not None:
    train_metrics = trainer.evaluate(train_dataset)
    for k, v in train_metrics.items():
        results[f"train_{k}"] = v
else:
    results["train_metrics"] = "No train dataset provided"

if val_dataset is not None:
    val_metrics = trainer.evaluate(val_dataset)
    for k, v in val_metrics.items():
        results[f"val_{k}"] = v
else:
    results["val_metrics"] = "No val dataset provided"

if test_dataset is not None:
    test_metrics = trainer.evaluate(test_dataset)
    for k, v in test_metrics.items():
        results[f"test_{k}"] = v
else:
    results["test_metrics"] = "No test dataset provided"

return results

# def log_experiment_results_json(experiment_meta, model_details, run_metrics,
# ↪ log_file):
#     """
#     Logs experiment metadata, model details, and metrics to a JSON lines file.
#     Automatically concatenates the 'checkpoint_path' to the 'model_lineage'.
#     """
#     checkpoint_path = model_details.get("checkpoint_path")
#     if checkpoint_path:
#         if "model_lineage" not in model_details:
#             model_details["model_lineage"] = ""
#         if model_details["model_lineage"]:
#             model_details["model_lineage"] += " -> "
#         model_details["model_lineage"] += checkpoint_path

#     record = {
#         "timestamp": str(datetime.datetime.now()),
#         "experiment_meta": experiment_meta,
#         "model_details": model_details,
#         "run_metrics": run_metrics
#     }

```

```

#     with open(log_file, "a", encoding="utf-8") as f:
#         json.dump(record, f)
#         f.write("\n")

def log_experiment_results_json(experiment_meta, model_details, run_metrics,
    ↪log_file):
    """
    Logs experiment metadata, model details, and metrics to a JSON lines file.
    Automatically concatenates the 'checkpoint_path' to the 'model_lineage'
    and uses Pacific time for the timestamp.
    """
    checkpoint_path = model_details.get("checkpoint_path")
    if checkpoint_path:
        if "model_lineage" not in model_details:
            model_details["model_lineage"] = ""
        if model_details["model_lineage"]:
            model_details["model_lineage"] += " -> "
        model_details["model_lineage"] += checkpoint_path

    pacific_time = datetime.now(zoneinfo.ZoneInfo("America/Los_Angeles")) # ↪
    ↪update to support pacific time
    timestamp_str = pacific_time.isoformat()

    record = {
        "timestamp": timestamp_str,
        "experiment_meta": experiment_meta,
        "model_details": model_details,
        "run_metrics": run_metrics
    }

    with open(log_file, "a", encoding="utf-8") as f:
        json.dump(record, f)
        f.write("\n")

```

## 0.2.2 Experiment Cohort Design

```

[ ]: # Define Experiment Parameters

named_model = "bert-base-cased"
# named_model = "roberta-base"
# named_model = "bert-large"
# named_model = "roberta-large"
# named_model = "" # modern bert

# learning_rate = 1e-3
# learning_rate = 1e-4
learning_rate = 1e-5

```

```

# learning_rate = 5e-6
# learning_rate = 5e-7
# learning_rate = 5e-8

# num_epochs = 1
# num_epochs = 3
# num_epochs = 5
num_epochs = 25
# num_epochs = 15
# num_epochs = 20

# length_max = 128
length_max = 256
# length_max = 348
# length_max = 512

# size_batch = 1
# size_batch = 4
# size_batch = 8
size_batch = 16
# size_batch = 24
# size_batch = 32
# size_batch = 64
# size_batch = 128

# regularization_weight_decay = 0
regularization_weight_decay = 0.1
# regularization_weight_decay = 0.5

y_col = "binary_complexity"
# y_col = "complexity"

x_task = "single"
# x_task = "multi"

# x_col = "sentence"
x_col = "sentence_no_contractions"
# x_col = "pos_sequence"
# x_col = "dep_sequence"
# x_col = "morph_sequence"

if x_task == "single":
    df_train = train_single_df
    df_val = train_val_single_df
    df_test = test_single_df
else:
    df_train = train_multi_df

```

```

df_val    = trial_val_multi_df
df_test   = test_multi_df

custom_config = BertConfig.from_pretrained("bert-base-cased")

custom_config.hidden_dropout_prob = 0.1
# custom_config.intermediate_size = 3072
# custom_config.intermediate_size = 6144
# custom_config.num_attention_heads = 12
# custom_config.num_hidden_layers = 12
custom_config.gradient_checkpointing = False
custom_config.attention_probs_dropout_prob = 0.1
# custom_config.max_position_embeddings = 512
# custom_config.type_vocab_size = 2
custom_config.hidden_act = "gelu" # alts: "relu" "silu"
# custom_config.vocab_size = 28996 # must match

# model.bert.pooler.activation = nn.ReLU() # Tanh() replaced as the pooler
↳ layer activation function in side-by-side with 1.1

```

```

[ ]: def train_transformer_model(
    model,
    tokenizer,
    train_dataset,
    val_dataset,
    output_dir=dir_results,
    num_epochs=num_epochs,
    batch_size=size_batch,
    lr=learning_rate,
    weight_decay=regularization_weight_decay
):
    """
    Sets up a Trainer and trains the model for 'num_epochs' using the given
    ↳ dataset.
    Returns the trained model and the Trainer object for possible re-use or
    ↳ analysis.
    """

    training_args = TrainingArguments(
        output_dir=output_dir,
        num_train_epochs=num_epochs,
        per_device_train_batch_size=batch_size,
        per_device_eval_batch_size=batch_size,
        evaluation_strategy="epoch",
        save_strategy="no",
        logging_strategy="epoch",
        learning_rate=lr,

```

```

        weight_decay=weight_decay,
        report_to=["none"], # or "wandb"
        warmup_steps=1
    )

    trainer = Trainer(
        model=model,
        args=training_args,
        train_dataset=train_dataset,
        eval_dataset=val_dataset,
        tokenizer=tokenizer, # optional
        compute_metrics=compute_metrics
    )

    trainer.train()
    return model, trainer

```

---



---

### Model Inspection \*\* Run \*\*

```

[ ]: print("model checkpoints:", dir_models)
     # !ls /content/drive/MyDrive/266-final/models/

```

model checkpoints: /content/drive/MyDrive/266-final/models/

```

[ ]: # Load Model & Tokenizer
     # model, tokenizer = get_model_and_tokenizer(named_model) # deprecated argument,
     # ↳ structure
     # model, tokenizer = get_model_and_tokenizer("/content/drive/MyDrive/266-final/
     # ↳ models/....") # proposed argument usage for checkpointed models

     # for name, param in model.named_parameters():
     #     print(name)

     model, tokenizer = get_model_and_tokenizer(
         remote_model_name="bert-base-cased",
         local_model_path=None,
         config=custom_config
     )

     # model, tokenizer = get_model_and_tokenizer(
     #     local_model_path="my_local_bert_path",
     #     config=custom_config
     # )

     print("=====")

```

```

print(named_model, ":")
print("=====")
# print(model)
print("=====")
print(model.config)
print("=====")
print("num_parameters:", model.num_parameters())
print("=====")
print("num_trainable_parameters:", model.num_parameters(only_trainable=True))

```

Loading from Hugging Face model: bert-base-cased

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-cased and are newly initialized:

['classifier.bias', 'classifier.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

=====

bert-base-cased :

=====

=====

```

BertConfig {
  "_attn_implementation_autoset": true,
  "architectures": [
    "BertForMaskedLM"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "gradient_checkpointing": false,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "pad_token_id": 0,
  "position_embedding_type": "absolute",
  "torch_dtype": "float32",
  "transformers_version": "4.50.3",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 28996
}

```

```

=====
num_parameters: 108311810
=====
num_trainable_parameters: 108311810

```

### Layer Configuration \*\* Run \*\*

```

[ ]: # Freeze/Unfreeze Layers & Additional Activation Function Configuration

layers_to_unfreeze = [
    # "bert.embeddings.",
    # "bert.encoder.layer.0.",
    # "bert.encoder.layer.1.",
    "bert.encoder.layer.8.",
    "bert.encoder.layer.9.",
    "bert.encoder.layer.10.",
    "bert.encoder.layer.11.",
    "bert.pooler.",
    "classifier.",
]

freeze_unfreeze_layers(model, layers_to_unfreeze=layers_to_unfreeze)

for name, param in model.named_parameters():
    print(name, "requires_grad=", param.requires_grad)

print("\nLayers that are 'True' are trainable. 'False' are frozen.")

print("=====")
print(named_model, ":")
print("=====")
# print(model)
print("=====")
print(model.config)
print("=====")
print("num_parameters:", model.num_parameters())
print("=====")
print("num_trainable_parameters:", model.num_parameters(only_trainable=True))

bert.embeddings.word_embeddings.weight requires_grad= False
bert.embeddings.position_embeddings.weight requires_grad= False
bert.embeddings.token_type_embeddings.weight requires_grad= False
bert.embeddings.LayerNorm.weight requires_grad= False
bert.embeddings.LayerNorm.bias requires_grad= False
bert.encoder.layer.0.attention.self.query.weight requires_grad= False
bert.encoder.layer.0.attention.self.query.bias requires_grad= False
bert.encoder.layer.0.attention.self.key.weight requires_grad= False
bert.encoder.layer.0.attention.self.key.bias requires_grad= False
bert.encoder.layer.0.attention.self.value.weight requires_grad= False

```



[illegible]



bert.encoder.layer.6.attention.self.value.bias requires\_grad= False  
bert.encoder.layer.6.attention.output.dense.weight requires\_grad= False  
bert.encoder.layer.6.attention.output.dense.bias requires\_grad= False  
bert.encoder.layer.6.attention.output.LayerNorm.weight requires\_grad= False  
bert.encoder.layer.6.attention.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.6.intermediate.dense.weight requires\_grad= False  
bert.encoder.layer.6.intermediate.dense.bias requires\_grad= False  
bert.encoder.layer.6.output.dense.weight requires\_grad= False  
bert.encoder.layer.6.output.dense.bias requires\_grad= False  
bert.encoder.layer.6.output.LayerNorm.weight requires\_grad= False  
bert.encoder.layer.6.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.7.attention.self.query.weight requires\_grad= False  
bert.encoder.layer.7.attention.self.query.bias requires\_grad= False  
bert.encoder.layer.7.attention.self.key.weight requires\_grad= False  
bert.encoder.layer.7.attention.self.key.bias requires\_grad= False  
bert.encoder.layer.7.attention.self.value.weight requires\_grad= False  
bert.encoder.layer.7.attention.self.value.bias requires\_grad= False  
bert.encoder.layer.7.attention.output.dense.weight requires\_grad= False  
bert.encoder.layer.7.attention.output.dense.bias requires\_grad= False  
bert.encoder.layer.7.attention.output.LayerNorm.weight requires\_grad= False  
bert.encoder.layer.7.attention.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.7.intermediate.dense.weight requires\_grad= False  
bert.encoder.layer.7.intermediate.dense.bias requires\_grad= False  
bert.encoder.layer.7.output.dense.weight requires\_grad= False  
bert.encoder.layer.7.output.dense.bias requires\_grad= False  
bert.encoder.layer.7.output.LayerNorm.weight requires\_grad= False  
bert.encoder.layer.7.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.8.attention.self.query.weight requires\_grad= True  
bert.encoder.layer.8.attention.self.query.bias requires\_grad= True  
bert.encoder.layer.8.attention.self.key.weight requires\_grad= True  
bert.encoder.layer.8.attention.self.key.bias requires\_grad= True  
bert.encoder.layer.8.attention.self.value.weight requires\_grad= True  
bert.encoder.layer.8.attention.self.value.bias requires\_grad= True  
bert.encoder.layer.8.attention.output.dense.weight requires\_grad= True  
bert.encoder.layer.8.attention.output.dense.bias requires\_grad= True  
bert.encoder.layer.8.attention.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.8.attention.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.8.intermediate.dense.weight requires\_grad= True  
bert.encoder.layer.8.intermediate.dense.bias requires\_grad= True  
bert.encoder.layer.8.output.dense.weight requires\_grad= True  
bert.encoder.layer.8.output.dense.bias requires\_grad= True  
bert.encoder.layer.8.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.8.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.9.attention.self.query.weight requires\_grad= True  
bert.encoder.layer.9.attention.self.query.bias requires\_grad= True  
bert.encoder.layer.9.attention.self.key.weight requires\_grad= True  
bert.encoder.layer.9.attention.self.key.bias requires\_grad= True  
bert.encoder.layer.9.attention.self.value.weight requires\_grad= True

bert.encoder.layer.9.attention.self.value.bias requires\_grad= True  
bert.encoder.layer.9.attention.output.dense.weight requires\_grad= True  
bert.encoder.layer.9.attention.output.dense.bias requires\_grad= True  
bert.encoder.layer.9.attention.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.9.attention.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.9.intermediate.dense.weight requires\_grad= True  
bert.encoder.layer.9.intermediate.dense.bias requires\_grad= True  
bert.encoder.layer.9.output.dense.weight requires\_grad= True  
bert.encoder.layer.9.output.dense.bias requires\_grad= True  
bert.encoder.layer.9.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.9.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.10.attention.self.query.weight requires\_grad= True  
bert.encoder.layer.10.attention.self.query.bias requires\_grad= True  
bert.encoder.layer.10.attention.self.key.weight requires\_grad= True  
bert.encoder.layer.10.attention.self.key.bias requires\_grad= True  
bert.encoder.layer.10.attention.self.value.weight requires\_grad= True  
bert.encoder.layer.10.attention.self.value.bias requires\_grad= True  
bert.encoder.layer.10.attention.output.dense.weight requires\_grad= True  
bert.encoder.layer.10.attention.output.dense.bias requires\_grad= True  
bert.encoder.layer.10.attention.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.10.attention.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.10.intermediate.dense.weight requires\_grad= True  
bert.encoder.layer.10.intermediate.dense.bias requires\_grad= True  
bert.encoder.layer.10.output.dense.weight requires\_grad= True  
bert.encoder.layer.10.output.dense.bias requires\_grad= True  
bert.encoder.layer.10.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.10.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.11.attention.self.query.weight requires\_grad= True  
bert.encoder.layer.11.attention.self.query.bias requires\_grad= True  
bert.encoder.layer.11.attention.self.key.weight requires\_grad= True  
bert.encoder.layer.11.attention.self.key.bias requires\_grad= True  
bert.encoder.layer.11.attention.self.value.weight requires\_grad= True  
bert.encoder.layer.11.attention.self.value.bias requires\_grad= True  
bert.encoder.layer.11.attention.output.dense.weight requires\_grad= True  
bert.encoder.layer.11.attention.output.dense.bias requires\_grad= True  
bert.encoder.layer.11.attention.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.11.attention.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.11.intermediate.dense.weight requires\_grad= True  
bert.encoder.layer.11.intermediate.dense.bias requires\_grad= True  
bert.encoder.layer.11.output.dense.weight requires\_grad= True  
bert.encoder.layer.11.output.dense.bias requires\_grad= True  
bert.encoder.layer.11.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.11.output.LayerNorm.bias requires\_grad= True  
bert.pooler.dense.weight requires\_grad= True  
bert.pooler.dense.bias requires\_grad= True  
classifier.weight requires\_grad= True  
classifier.bias requires\_grad= True

Layers that are 'True' are trainable. 'False' are frozen.

```
=====
bert-base-cased :
=====
BertConfig {
  "_attn_implementation_autoset": true,
  "architectures": [
    "BertForMaskedLM"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "gradient_checkpointing": false,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 768,
  "initializer_range": 0.02,
  "intermediate_size": 3072,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "pad_token_id": 0,
  "position_embedding_type": "absolute",
  "torch_dtype": "float32",
  "transformers_version": "4.50.3",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 28996
}
```

```
=====
num_parameters: 108311810
=====
num_trainable_parameters: 28943618
```

### Dataset Preparation \*\* Run \*\*

```
[ ]: # Tokenize & Prepare Datasets

train_data_hf = prepare_dataset(
    df_train,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max
)
```





```

y_col = "binary_complexity"
# y_col = "complexity"
#####
# x_task = "single"
x_task = "multi"
if x_task == "single":
    df_train = train_single_df
    df_val = trial_val_single_df
    df_test = test_single_df
else:
    df_train = train_multi_df
    df_val = trial_val_multi_df
    df_test = test_multi_df
#####
# Tokenize & Prepare Datasets
train_data_hf = prepare_dataset(
    df_train,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
val_data_hf = prepare_dataset(
    df_val,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
test_data_hf = prepare_dataset(
    df_test,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
print("Datasets prepared. Sample from train_data_hf:\n", train_data_hf[10])
# print("Datasets prepared. Sample from train_data_hf:\n", val_data_hf[10])
# print("Datasets prepared. Sample from train_data_hf:\n", test_data_hf[10])
#####
# custom_config = BertConfig.from_pretrained("bert-base-cased")
# custom_config.hidden_act = "gelu" # alts: "relu" "silu"
# custom_config.attention_probs_dropout_prob = 0.1
# custom_config.hidden_dropout_prob = 0.1
# custom_config.gradient_checkpointing = False
#####
model, tokenizer = get_model_and_tokenizer(
    remote_model_name="bert-base-cased",
    local_model_path=None,
    config=None)

```



```

#####
# model, tokenizer = get_model_and_tokenizer(
#     remote_model_name=None
#     local_model_path="...CONFIGURE_PATH...",
#     config=custom_config)
print("=====")
print(named_model, ":")
print("=====")
print("num_parameters:", model.num_parameters())
print("num_trainable_parameters at load:", model.
    ↪ num_parameters(only_trainable=True))
print("=====")
print("model lineage:", MODEL_LINEAGE)
print("=====")
#####
layers_to_unfreeze = [
    # "bert.embeddings.",
    # "bert.encoder.layer.0.",
    # "bert.encoder.layer.1.",
    # "bert.encoder.layer.8.",
    # "bert.encoder.layer.9.",
    # "bert.encoder.layer.10.",
    "bert.encoder.layer.11.",
    "bert.pooler.",
    "classifier.",
]
freeze_unfreeze_layers(model, layers_to_unfreeze=layers_to_unfreeze)
print(model.config)
print("=====")
print("num_parameters:", model.num_parameters())
print("num_trainable_parameters:", model.num_parameters(only_trainable=True))
print("=====")
#####
print("Experiment configuration used with this experiment:")
print("model used:", named_model)
print("learning rate used:", learning_rate)
print("number of epochs:", num_epochs)
print("maximum sequence length:", length_max)
print("batch size used:", size_batch)
print("regularization value:", regularization_weight_decay)
print("outcome variable:", y_col)
print("task:", x_task)
print("input column:", x_col)

```

Map: 0%| | 0/1300 [00:00<?, ? examples/s]

Map: 0%| | 0/250 [00:00<?, ? examples/s]

Map: 0%| | 0/250 [00:00<?, ? examples/s]

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-cased and are newly initialized:

Datasets prepared. Sample from train\_data\_hf:

[illegible]

Loading from Hugging Face model: bert-base-cased

```
num_parameters: 108311810
```

```
model lineage: {'type': 'huggingface_hub', 'path': 'bert-base-cased',
'timestamp': '2025-04-14 01:53:36'}
```

```
BertConfig {
```

1.

```

    "hidden_size": 768,
    "initializer_range": 0.02,
    "intermediate_size": 3072,
    "layer_norm_eps": 1e-12,
    "max_position_embeddings": 512,
    "model_type": "bert",
    "num_attention_heads": 12,
    "num_hidden_layers": 12,
    "pad_token_id": 0,
    "position_embedding_type": "absolute",
    "torch_dtype": "float32",
    "transformers_version": "4.50.3",
    "type_vocab_size": 2,
    "use_cache": true,
    "vocab_size": 28996
}

```

```

=====

```

```

num_parameters: 108311810
num_trainable_parameters: 7680002
=====

```

Experiment configuration used with this experiment:

```

model used: bert-base-cased
learning rate used: 5e-06
number of epochs: 1
maximum sequence length: 128
batch size used: 128
regularization value: 0.5
outcome variable: binary_complexity
task: multi
input column: sentence_no_contractions

```

```

[ ]: # Train & Evaluate
trained_model, trainer_obj = train_transformer_model(
    model = model,
    tokenizer = tokenizer,
    train_dataset = train_data_hf,
    val_dataset = val_data_hf,
    output_dir = dir_results,
    num_epochs = num_epochs,
    batch_size = size_batch,
    lr = learning_rate,
    weight_decay = regularization_weight_decay)
metrics = trainer_obj.evaluate()
print("Validation metrics:", metrics)
test_metrics = trainer_obj.evaluate(test_data_hf) if test_data_hf else None
print("Test metrics:", test_metrics)

```

```

/usr/local/lib/python3.11/dist-packages/transformers/training_args.py:1611:
FutureWarning: `evaluation_strategy` is deprecated and will be removed in
version 4.46 of Transformers. Use `eval_strategy` instead
    warnings.warn(
<ipython-input-20-81222a87c90b>:31: FutureWarning: `tokenizer` is deprecated and
will be removed in version 5.0.0 for `Trainer.__init__`. Use `processing_class`
instead.
    trainer = Trainer(

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Validation metrics: {'eval_loss': 0.71732497215271, 'eval_accuracy': 0.436,
'eval_precision': 0.43373493975903615, 'eval_recall': 1.0, 'eval_f1':
0.6050420168067226, 'eval_runtime': 1.4688, 'eval_samples_per_second': 170.205,
'eval_steps_per_second': 1.362, 'epoch': 1.0}
Test metrics: {'eval_loss': 0.6920300722122192, 'eval_accuracy': 0.508,
'eval_precision': 0.508, 'eval_recall': 1.0, 'eval_f1': 0.6737400530503979,
'eval_runtime': 1.8379, 'eval_samples_per_second': 136.027,
'eval_steps_per_second': 1.088, 'epoch': 1.0}

```

```

[ ]: # save model checkpoint
# timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
pacific_time = datetime.now(zoneinfo.ZoneInfo("America/Los_Angeles"))
timestamp = pacific_time.isoformat()
model_save_path = os.path.join(dir_models,
    f"{x_task}_{named_model}_{y_col}_{timestamp}")
trainer_obj.save_model(model_save_path)
print(f"Model checkpoint saved to: {model_save_path}")
# log experiment results
experiment_info = {
    "model_name": named_model,
    "learning_rate": learning_rate,
    "epochs": num_epochs,
    "batch_size": size_batch,
    "weight_decay": regularization_weight_decay,
    "x_task": x_task,
    "x_col": x_col,
    "y_col": y_col,
    "layers_to_unfreeze": layers_to_unfreeze}
model_info = gather_model_details(trained_model)
all_run_metrics = gather_all_run_metrics(
    trainer=trainer_obj,
    train_dataset=train_data_hf,
    val_dataset=val_data_hf,
    test_dataset=test_data_hf)
log_experiment_results_json(
    experiment_meta=experiment_info,

```

```

    model_details=model_info,
    run_metrics=all_run_metrics,
    log_file=log_filepath)
print(f"EXPERIMENT LOGGED TO: {log_filepath}")

```

Model checkpoint saved to: /content/drive/MyDrive/266-final/models/multi\_bert-base-cased\_binary\_complexity\_2025-04-13T18:53:44.827658-07:00

<IPython.core.display.HTML object>

EXPERIMENT LOGGED TO:

/content/drive/MyDrive/266-final/results/experiment\_runs.txt

```

[ ]: prediction_output = trainer_obj.predict(test_data_hf)
raw_predictions = prediction_output.predictions
true_labels = prediction_output.label_ids
preds = np.argmax(raw_predictions, axis=1)
mismatch_indices = np.where(preds != true_labels)[0]

error_rows = []
for idx in mismatch_indices:
    text_value = df_test.iloc[idx][x_col]
    true_label_val = true_labels[idx]
    pred_label_val = preds[idx]

    error_rows.append({
        "hf_index": idx,
        "text": text_value,
        "true_label": true_label_val,
        "predicted_label": pred_label_val
    })

error_df = pd.DataFrame(error_rows)
df_test_for_merge = df_test.copy()
df_test_for_merge["error_matching_prefix"] = df_test_for_merge[x_col].str[:50]
# df_test_for_merge.drop(columns=[x_col], inplace=True)

error_df["error_matching_prefix"] = error_df["text"].str[:50]
error_df = error_df.merge(
    df_test_for_merge,
    on="error_matching_prefix",
    how="left",
    suffixes=("", "_source"))

error_df.to_csv("bert-base-cased_mismatches.csv", index=False)

# print("Number of misclassified samples:", len(error_df))
print("\nMerged error_df with extra columns:")

```

```

# display(error_df.head(15))

# print("\nConfusion Matrix:")
# cm = confusion_matrix(true_labels, preds)
# plt.figure(figsize=(8, 6))
# sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=True,
#             xticklabels=["Predicted 0", "Predicted 1"],
#             yticklabels=["True 0", "True 1"])
# plt.xlabel('Predicted Label')
# plt.ylabel('True Label')
# plt.title('Confusion Matrix')
# plt.tight_layout()
# plt.show()
# print("confusion matrix metrics: \n", cm)
# error_save_path = os.path.join(dir_results,
#     ↳ f"{x_task}_{named_model}_{y_col}_{x_col}_errors.csv")
# error_df.to_csv(error_save_path, index=False)
# print("Result saved to results directory.")
prediction_output = trainer_obj.predict(test_data_hf)
raw_predictions = prediction_output.predictions
true_labels = prediction_output.label_ids
preds = np.argmax(raw_predictions, axis=1)
df_test["true_label"] = true_labels
df_test["predicted_label"] = preds
df_test["is_incorrect"] = (df_test["predicted_label"] != df_test["true_label"])
df_test["avg_embedding"] = None
device = next(model.parameters()).device
for i in range(len(df_test)):
    text_value = df_test.iloc[i][x_col]
    e = tokenizer(text_value, return_tensors="pt", truncation=True,
        ↳ max_length=512).to(device)
    with torch.no_grad():
        emb = model.bert.embeddings.word_embeddings(e["input_ids"]).mean(dim=1).
        ↳ squeeze().cpu().numpy()
    df_test.at[i, "avg_embedding"] = emb
cm = confusion_matrix(df_test["true_label"], df_test["predicted_label"])
plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=True,
    ↳ xticklabels=["Predicted 0", "Predicted 1"], yticklabels=["True 0", "True 1"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.tight_layout()
plt.show()
# print("confusion matrix metrics:\n", cm)

df_plot = df_test.dropna(subset=["avg_embedding"]).copy()

```

```

embeddings = np.stack(df_plot["avg_embedding"].values)
pca = PCA(n_components=3)
reduced = pca.fit_transform(embeddings)
df_plot["pca_x"] = reduced[:,0]
df_plot["pca_y"] = reduced[:,1]
df_plot["pca_z"] = reduced[:,2]
colors = {
    ("bible", True): "red",
    ("bible", False): "orange",
    ("europarl", True): "yellow",
    ("europarl", False): "green",
    ("biomed", True): "blue",
    ("biomed", False): "purple"}
fig = plt.figure(figsize=(13,13))
ax = fig.add_subplot(111, projection='3d')
fig.set_facecolor("white")
ax.set_facecolor("white")
ax.xaxis._axinfo["grid"]["color"] = "gray"
ax.yaxis._axinfo["grid"]["color"] = "gray"
ax.zaxis._axinfo["grid"]["color"] = "gray"
for (corp, incorr), color in colors.items():
    subset = df_plot[(df_plot["corpus"]==corp) &
    ↪(df_plot["is_incorrect"]==incorr)]
    ax.scatter(subset["pca_x"], subset["pca_y"], subset["pca_z"], c=color,
    ↪s=20, alpha=0.8, label=f"{corp}, incorrect={incorr}")
ax.set_title("Average Embedding Value of Predictions, by Corpus", color="black")
ax.set_xlabel("PC1", color="black")
ax.set_ylabel("PC2", color="black")
ax.set_zlabel("PC3", color="black")
ax.legend(loc="best")
plt.show()
if "corpus" in df_test.columns:
    freqs = df_test["corpus"].value_counts()
    print("\nFrequency counts of corpus:", freqs)
    err_df = df_test[df_test["is_incorrect"]==True]
    corr_df = df_test[df_test["is_incorrect"]==False]
    err_counts = err_df["corpus"].value_counts()
    corr_counts = corr_df["corpus"].value_counts()
    print("\nCounts of corpus in misclassified:", err_counts)
    print("\nCounts of corpus in correctly classified:", corr_counts)
    print("\nProportions of corpus in misclassified:", err_counts/err_counts.
    ↪sum())
    print("\nProportions of corpus in correctly classified:", corr_counts/
    ↪corr_counts.sum())
    grouped_all = df_test.groupby("corpus")["avg_embedding"].apply(lambda x: np.
    ↪mean(np.stack(x.values), axis=0))

```

```

    grouped_err = err_df.groupby("corpus")["avg_embedding"].apply(lambda x: np.
↳mean(np.stack(x.values), axis=0))
    grouped_corr = corr_df.groupby("corpus")["avg_embedding"].apply(lambda x:
↳np.mean(np.stack(x.values), axis=0))
    # print("\nAvg embedding of each subcorpus overall:", grouped_all)
    # print("\nAvg embedding of each subcorpus misclassified:", grouped_err)
    # print("\nAvg embedding of each subcorpus correctly classified:",
↳grouped_corr)
err_stack = np.stack(df_test[df_test["is_incorrect"]==True]["avg_embedding"].
↳values)
corr_stack = np.stack(df_test[df_test["is_incorrect"]==False]["avg_embedding"].
↳values)
# print("\nAvg embedding of records predicted incorrectly:", err_stack.
↳mean(axis=0))
# print("Avg embedding of records predicted correctly:", corr_stack.
↳mean(axis=0))
error_df = df_test[df_test["is_incorrect"]==True].copy()
error_df.to_csv("misclassified_with_all_columns.csv", index=False)
error_save_path = os.path.join(dir_results,
↳f"{x_task}_{named_model}_{y_col}_{x_col}_errors.csv")
df_test.to_csv(error_save_path, index=False)
for corp in df_test["corpus"].unique():
    subset = df_test[df_test["corpus"]==corp]
    emb_true = subset[subset["is_incorrect"]==False]["avg_embedding"]
    emb_false = subset[subset["is_incorrect"]==True]["avg_embedding"]
    if len(emb_true)==0 or len(emb_false)==0:
        print(f"No valid data for subcorpus '{corp}'")
        continue
    p = np.mean(np.stack(emb_true.values), axis=0)
    q = np.mean(np.stack(emb_false.values), axis=0)
    p_exp = np.exp(p - np.max(p))
    q_exp = np.exp(q - np.max(q))
    p_sum = p_exp.sum()
    q_sum = q_exp.sum()
    if p_sum<=0 or q_sum<=0:
        print(f"Cannot form valid distributions for subcorpus '{corp}'")
        continue
    p_dist = p_exp / p_sum
    q_dist = q_exp / q_sum
    kl_pq = entropy(p_dist, q_dist)
    kl_qp = entropy(q_dist, p_dist)
    kl_sym = 0.5*(kl_pq + kl_qp)
    print(f"Subcorpus '{corp}' symmetric KL divergence: {kl_sym}")
error_save_path = os.path.join(dir_results,
↳f"{x_task}_{named_model}_{y_col}_{x_col}_errors.csv")
df_test.to_csv(error_save_path, index=False)

```





```
print("model results saved")
```

<IPython.core.display.HTML object>

Merged error\_df with extra columns:

<IPython.core.display.HTML object>





Frequency counts of corpus: corpus

bible	88
europarl	87
biomed	75

Name: count, dtype: int64

Counts of corpus in misclassified: corpus

bible	51
europarl	44
biomed	28

Name: count, dtype: int64

Counts of corpus in correctly classified: corpus

biomed 47

europarl 43

bible 37

Name: count, dtype: int64

Proportions of corpus in misclassified: corpus

bible 0.414634

europarl 0.357724

biomed 0.227642

Name: count, dtype: float64

Proportions of corpus in correctly classified: corpus

biomed 0.370079

europarl 0.338583

bible 0.291339

Name: count, dtype: float64

Subcorpus 'bible' symmetric KL divergence: 2.95432516672069e-06

Subcorpus 'biomed' symmetric KL divergence: 1.6558011492547724e-06

Subcorpus 'europarl' symmetric KL divergence: 3.947562264288915e-06

model results saved

## Result

### 0.2.4 bert-large-cased

```
[ ]: # Define Experiment Parameters
# named_model = "bert-base-cased"
# named_model = "roberta-base"
named_model = "bert-large-cased"
# named_model = "roberta-large"
# named_model = "" # modern bert
#####
regularization_weight_decay = 0.5
learning_rate = 5e-6
size_batch = 128
length_max = 128
num_epochs = 1
#####
# x_col = "sentence"
x_col = "sentence_no_contractions"
# x_col = "pos_sequence"
# x_col = "dep_sequence"
# x_col = "morph_sequence"
# x_col = "snc_pos_seq"
# x_col = "snc_pos_alt"
# x_col = "snc_morph_seq"
```

```

# x_col = "snc_morph_alt"
# x_col = "snc_dep_seq"
# x_col = "snc_dep_alt"
# x_col = "snc_morph_complexity_value"
#####
y_col = "binary_complexity"
# y_col = "complexity"
#####
# x_task = "single"
x_task = "multi"
if x_task == "single":
    df_train = train_single_df
    df_val = trial_val_single_df
    df_test = test_single_df
else:
    df_train = train_multi_df
    df_val = trial_val_multi_df
    df_test = test_multi_df
#####
# Tokenize & Prepare Datasets
train_data_hf = prepare_dataset(
    df_train,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
val_data_hf = prepare_dataset(
    df_val,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
test_data_hf = prepare_dataset(
    df_test,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
print("Datasets prepared. Sample from train_data_hf:\n", train_data_hf[10])
# print("Datasets prepared. Sample from train_data_hf:\n", val_data_hf[10])
# print("Datasets prepared. Sample from train_data_hf:\n", test_data_hf[10])
#####
# custom_config = BertConfig.from_pretrained("roberta-base")
# custom_config.hidden_act = "gelu" # alts: "relu" "silu"
# custom_config.attention_probs_dropout_prob = 0.1
# custom_config.hidden_dropout_prob = 0.1
# custom_config.gradient_checkpointing = False

```



Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-large-cased and are newly initialized:  
['classifier.bias', 'classifier.weight']  
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
=====
bert-large-cased :
=====
num_parameters: 333581314
num_trainable_parameters at load: 333581314
=====
model lineage: {'type': 'huggingface_hub', 'path': 'bert-large-cased',
'timestamp': '2025-04-14 01:54:01'}
=====
```

```
[ ]: print(model)
```

```
BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(28996, 1024, padding_idx=0)
      (position_embeddings): Embedding(512, 1024)
      (token_type_embeddings): Embedding(2, 1024)
      (LayerNorm): LayerNorm((1024,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-23): 24 x BertLayer(
          (attention): BertAttention(
            (self): BertSdpaSelfAttention(
              (query): Linear(in_features=1024, out_features=1024, bias=True)
              (key): Linear(in_features=1024, out_features=1024, bias=True)
              (value): Linear(in_features=1024, out_features=1024, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=1024, out_features=1024, bias=True)
              (LayerNorm): LayerNorm((1024,), eps=1e-12,
elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
        )
      )
      (intermediate): BertIntermediate(
        (dense): Linear(in_features=1024, out_features=4096, bias=True)
        (intermediate_act_fn): GELUActivation()
      )
      (output): BertOutput(
```

```

        (dense): Linear(in_features=4096, out_features=1024, bias=True)
        (LayerNorm): LayerNorm((1024,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
)
)
(pooler): BertPooler(
  (dense): Linear(in_features=1024, out_features=1024, bias=True)
  (activation): Tanh()
)
)
(dropout): Dropout(p=0.1, inplace=False)
(classifier): Linear(in_features=1024, out_features=2, bias=True)
)

```

```

[ ]: for name, param in model.named_parameters():
      print(name, "requires_grad=", param.requires_grad)

```

```

bert.embeddings.word_embeddings.weight requires_grad= True
bert.embeddings.position_embeddings.weight requires_grad= True
bert.embeddings.token_type_embeddings.weight requires_grad= True
bert.embeddings.LayerNorm.weight requires_grad= True
bert.embeddings.LayerNorm.bias requires_grad= True
bert.encoder.layer.0.attention.self.query.weight requires_grad= True
bert.encoder.layer.0.attention.self.query.bias requires_grad= True
bert.encoder.layer.0.attention.self.key.weight requires_grad= True
bert.encoder.layer.0.attention.self.key.bias requires_grad= True
bert.encoder.layer.0.attention.self.value.weight requires_grad= True
bert.encoder.layer.0.attention.self.value.bias requires_grad= True
bert.encoder.layer.0.attention.output.dense.weight requires_grad= True
bert.encoder.layer.0.attention.output.dense.bias requires_grad= True
bert.encoder.layer.0.attention.output.LayerNorm.weight requires_grad= True
bert.encoder.layer.0.attention.output.LayerNorm.bias requires_grad= True
bert.encoder.layer.0.intermediate.dense.weight requires_grad= True
bert.encoder.layer.0.intermediate.dense.bias requires_grad= True
bert.encoder.layer.0.output.dense.weight requires_grad= True
bert.encoder.layer.0.output.dense.bias requires_grad= True
bert.encoder.layer.0.output.LayerNorm.weight requires_grad= True
bert.encoder.layer.0.output.LayerNorm.bias requires_grad= True
bert.encoder.layer.1.attention.self.query.weight requires_grad= True
bert.encoder.layer.1.attention.self.query.bias requires_grad= True
bert.encoder.layer.1.attention.self.key.weight requires_grad= True
bert.encoder.layer.1.attention.self.key.bias requires_grad= True
bert.encoder.layer.1.attention.self.value.weight requires_grad= True
bert.encoder.layer.1.attention.self.value.bias requires_grad= True
bert.encoder.layer.1.attention.output.dense.weight requires_grad= True
bert.encoder.layer.1.attention.output.dense.bias requires_grad= True

```







bert.encoder.layer.7.attention.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.7.attention.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.7.intermediate.dense.weight requires\_grad= True  
bert.encoder.layer.7.intermediate.dense.bias requires\_grad= True  
bert.encoder.layer.7.output.dense.weight requires\_grad= True  
bert.encoder.layer.7.output.dense.bias requires\_grad= True  
bert.encoder.layer.7.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.7.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.8.attention.self.query.weight requires\_grad= True  
bert.encoder.layer.8.attention.self.query.bias requires\_grad= True  
bert.encoder.layer.8.attention.self.key.weight requires\_grad= True  
bert.encoder.layer.8.attention.self.key.bias requires\_grad= True  
bert.encoder.layer.8.attention.self.value.weight requires\_grad= True  
bert.encoder.layer.8.attention.self.value.bias requires\_grad= True  
bert.encoder.layer.8.attention.output.dense.weight requires\_grad= True  
bert.encoder.layer.8.attention.output.dense.bias requires\_grad= True  
bert.encoder.layer.8.attention.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.8.attention.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.8.intermediate.dense.weight requires\_grad= True  
bert.encoder.layer.8.intermediate.dense.bias requires\_grad= True  
bert.encoder.layer.8.output.dense.weight requires\_grad= True  
bert.encoder.layer.8.output.dense.bias requires\_grad= True  
bert.encoder.layer.8.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.8.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.9.attention.self.query.weight requires\_grad= True  
bert.encoder.layer.9.attention.self.query.bias requires\_grad= True  
bert.encoder.layer.9.attention.self.key.weight requires\_grad= True  
bert.encoder.layer.9.attention.self.key.bias requires\_grad= True  
bert.encoder.layer.9.attention.self.value.weight requires\_grad= True  
bert.encoder.layer.9.attention.self.value.bias requires\_grad= True  
bert.encoder.layer.9.attention.output.dense.weight requires\_grad= True  
bert.encoder.layer.9.attention.output.dense.bias requires\_grad= True  
bert.encoder.layer.9.attention.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.9.attention.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.9.intermediate.dense.weight requires\_grad= True  
bert.encoder.layer.9.intermediate.dense.bias requires\_grad= True  
bert.encoder.layer.9.output.dense.weight requires\_grad= True  
bert.encoder.layer.9.output.dense.bias requires\_grad= True  
bert.encoder.layer.9.output.LayerNorm.weight requires\_grad= True  
bert.encoder.layer.9.output.LayerNorm.bias requires\_grad= True  
bert.encoder.layer.10.attention.self.query.weight requires\_grad= True  
bert.encoder.layer.10.attention.self.query.bias requires\_grad= True  
bert.encoder.layer.10.attention.self.key.weight requires\_grad= True  
bert.encoder.layer.10.attention.self.key.bias requires\_grad= True  
bert.encoder.layer.10.attention.self.value.weight requires\_grad= True  
bert.encoder.layer.10.attention.self.value.bias requires\_grad= True  
bert.encoder.layer.10.attention.output.dense.weight requires\_grad= True  
bert.encoder.layer.10.attention.output.dense.bias requires\_grad= True

[illegible]

[illegible]

[illegible]

[illegible]

```

bert.encoder.layer.22.attention.output.LayerNorm.weight requires_grad= True
bert.encoder.layer.22.attention.output.LayerNorm.bias requires_grad= True
bert.encoder.layer.22.intermediate.dense.weight requires_grad= True
bert.encoder.layer.22.intermediate.dense.bias requires_grad= True
bert.encoder.layer.22.output.dense.weight requires_grad= True
bert.encoder.layer.22.output.dense.bias requires_grad= True
bert.encoder.layer.22.output.LayerNorm.weight requires_grad= True
bert.encoder.layer.22.output.LayerNorm.bias requires_grad= True
bert.encoder.layer.23.attention.self.query.weight requires_grad= True
bert.encoder.layer.23.attention.self.query.bias requires_grad= True
bert.encoder.layer.23.attention.self.key.weight requires_grad= True
bert.encoder.layer.23.attention.self.key.bias requires_grad= True
bert.encoder.layer.23.attention.self.value.weight requires_grad= True
bert.encoder.layer.23.attention.self.value.bias requires_grad= True
bert.encoder.layer.23.attention.output.dense.weight requires_grad= True
bert.encoder.layer.23.attention.output.dense.bias requires_grad= True
bert.encoder.layer.23.attention.output.LayerNorm.weight requires_grad= True
bert.encoder.layer.23.attention.output.LayerNorm.bias requires_grad= True
bert.encoder.layer.23.intermediate.dense.weight requires_grad= True
bert.encoder.layer.23.intermediate.dense.bias requires_grad= True
bert.encoder.layer.23.output.dense.weight requires_grad= True
bert.encoder.layer.23.output.dense.bias requires_grad= True
bert.encoder.layer.23.output.LayerNorm.weight requires_grad= True
bert.encoder.layer.23.output.LayerNorm.bias requires_grad= True
bert.pooler.dense.weight requires_grad= True
bert.pooler.dense.bias requires_grad= True
classifier.weight requires_grad= True
classifier.bias requires_grad= True

```

```

[ ]: #####
layers_to_unfreeze = [
    "bert.encoder.layer.23.",
    "bert.pooler.",
    "classifier.",
]
freeze_unfreeze_layers(model, layers_to_unfreeze=layers_to_unfreeze)
print(model.config)
print("=====")
print("num_parameters:", model.num_parameters())
print("num_trainable_parameters:", model.num_parameters(only_trainable=True))
print("=====")
#####
print("Experiment configuration used with this experiment:")
print("model used:", named_model)
print("learning rate used:", learning_rate)
print("number of epochs:", num_epochs)
print("maximum sequence length:", length_max)

```

```

print("batch size used:", size_batch)
print("regularization value:", regularization_weight_decay)
print("outcome variable:", y_col)
print("task:", x_task)
print("input column:", x_col)
#####
print("=====")
print("num_trainable_parameters:", model.num_parameters(only_trainable=True))

```

```

BertConfig {
  "_attn_implementation_autoset": true,
  "architectures": [
    "BertForMaskedLM"
  ],
  "attention_probs_dropout_prob": 0.1,
  "classifier_dropout": null,
  "directionality": "bidi",
  "gradient_checkpointing": false,
  "hidden_act": "gelu",
  "hidden_dropout_prob": 0.1,
  "hidden_size": 1024,
  "initializer_range": 0.02,
  "intermediate_size": 4096,
  "layer_norm_eps": 1e-12,
  "max_position_embeddings": 512,
  "model_type": "bert",
  "num_attention_heads": 16,
  "num_hidden_layers": 24,
  "pad_token_id": 0,
  "pooler_fc_size": 768,
  "pooler_num_attention_heads": 12,
  "pooler_num_fc_layers": 3,
  "pooler_size_per_head": 128,
  "pooler_type": "first_token_transform",
  "position_embedding_type": "absolute",
  "torch_dtype": "float32",
  "transformers_version": "4.50.3",
  "type_vocab_size": 2,
  "use_cache": true,
  "vocab_size": 28996
}

```

=====

```

num_parameters: 333581314
num_trainable_parameters: 13647874

```

=====

Experiment configuration used with this experiment:  
model used: bert-large-cased



```

learning rate used: 5e-06
number of epochs: 1
maximum sequence length: 128
batch size used: 128
regularization value: 0.5
outcome variable: binary_complexity
task: multi
input column: sentence_no_contractions
=====
num_trainable_parameters: 13647874

```

```
[ ]: model.resize_token_embeddings(len(tokenizer))
```

```
[ ]: Embedding(28996, 1024, padding_idx=0)
```

```
[ ]: for name, param in model.named_parameters():
      print(name, "requires_grad=", param.requires_grad)
```

```

bert.embeddings.word_embeddings.weight requires_grad= False
bert.embeddings.position_embeddings.weight requires_grad= False
bert.embeddings.token_type_embeddings.weight requires_grad= False
bert.embeddings.LayerNorm.weight requires_grad= False
bert.embeddings.LayerNorm.bias requires_grad= False
bert.encoder.layer.0.attention.self.query.weight requires_grad= False
bert.encoder.layer.0.attention.self.query.bias requires_grad= False
bert.encoder.layer.0.attention.self.key.weight requires_grad= False
bert.encoder.layer.0.attention.self.key.bias requires_grad= False
bert.encoder.layer.0.attention.self.value.weight requires_grad= False
bert.encoder.layer.0.attention.self.value.bias requires_grad= False
bert.encoder.layer.0.attention.output.dense.weight requires_grad= False
bert.encoder.layer.0.attention.output.dense.bias requires_grad= False
bert.encoder.layer.0.attention.output.LayerNorm.weight requires_grad= False
bert.encoder.layer.0.attention.output.LayerNorm.bias requires_grad= False
bert.encoder.layer.0.intermediate.dense.weight requires_grad= False
bert.encoder.layer.0.intermediate.dense.bias requires_grad= False
bert.encoder.layer.0.output.dense.weight requires_grad= False
bert.encoder.layer.0.output.dense.bias requires_grad= False
bert.encoder.layer.0.output.LayerNorm.weight requires_grad= False
bert.encoder.layer.0.output.LayerNorm.bias requires_grad= False
bert.encoder.layer.1.attention.self.query.weight requires_grad= False
bert.encoder.layer.1.attention.self.query.bias requires_grad= False
bert.encoder.layer.1.attention.self.key.weight requires_grad= False
bert.encoder.layer.1.attention.self.key.bias requires_grad= False
bert.encoder.layer.1.attention.self.value.weight requires_grad= False
bert.encoder.layer.1.attention.self.value.bias requires_grad= False
bert.encoder.layer.1.attention.output.dense.weight requires_grad= False
bert.encoder.layer.1.attention.output.dense.bias requires_grad= False
bert.encoder.layer.1.attention.output.LayerNorm.weight requires_grad= False

```

[illegible]



bert.encoder.layer.7.attention.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.7.intermediate.dense.weight requires\_grad= False  
bert.encoder.layer.7.intermediate.dense.bias requires\_grad= False  
bert.encoder.layer.7.output.dense.weight requires\_grad= False  
bert.encoder.layer.7.output.dense.bias requires\_grad= False  
bert.encoder.layer.7.output.LayerNorm.weight requires\_grad= False  
bert.encoder.layer.7.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.8.attention.self.query.weight requires\_grad= False  
bert.encoder.layer.8.attention.self.query.bias requires\_grad= False  
bert.encoder.layer.8.attention.self.key.weight requires\_grad= False  
bert.encoder.layer.8.attention.self.key.bias requires\_grad= False  
bert.encoder.layer.8.attention.self.value.weight requires\_grad= False  
bert.encoder.layer.8.attention.self.value.bias requires\_grad= False  
bert.encoder.layer.8.attention.output.dense.weight requires\_grad= False  
bert.encoder.layer.8.attention.output.dense.bias requires\_grad= False  
bert.encoder.layer.8.attention.output.LayerNorm.weight requires\_grad= False  
bert.encoder.layer.8.attention.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.8.intermediate.dense.weight requires\_grad= False  
bert.encoder.layer.8.intermediate.dense.bias requires\_grad= False  
bert.encoder.layer.8.output.dense.weight requires\_grad= False  
bert.encoder.layer.8.output.dense.bias requires\_grad= False  
bert.encoder.layer.8.output.LayerNorm.weight requires\_grad= False  
bert.encoder.layer.8.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.9.attention.self.query.weight requires\_grad= False  
bert.encoder.layer.9.attention.self.query.bias requires\_grad= False  
bert.encoder.layer.9.attention.self.key.weight requires\_grad= False  
bert.encoder.layer.9.attention.self.key.bias requires\_grad= False  
bert.encoder.layer.9.attention.self.value.weight requires\_grad= False  
bert.encoder.layer.9.attention.self.value.bias requires\_grad= False  
bert.encoder.layer.9.attention.output.dense.weight requires\_grad= False  
bert.encoder.layer.9.attention.output.dense.bias requires\_grad= False  
bert.encoder.layer.9.attention.output.LayerNorm.weight requires\_grad= False  
bert.encoder.layer.9.attention.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.9.intermediate.dense.weight requires\_grad= False  
bert.encoder.layer.9.intermediate.dense.bias requires\_grad= False  
bert.encoder.layer.9.output.dense.weight requires\_grad= False  
bert.encoder.layer.9.output.dense.bias requires\_grad= False  
bert.encoder.layer.9.output.LayerNorm.weight requires\_grad= False  
bert.encoder.layer.9.output.LayerNorm.bias requires\_grad= False  
bert.encoder.layer.10.attention.self.query.weight requires\_grad= False  
bert.encoder.layer.10.attention.self.query.bias requires\_grad= False  
bert.encoder.layer.10.attention.self.key.weight requires\_grad= False  
bert.encoder.layer.10.attention.self.key.bias requires\_grad= False  
bert.encoder.layer.10.attention.self.value.weight requires\_grad= False  
bert.encoder.layer.10.attention.self.value.bias requires\_grad= False  
bert.encoder.layer.10.attention.output.dense.weight requires\_grad= False  
bert.encoder.layer.10.attention.output.dense.bias requires\_grad= False  
bert.encoder.layer.10.attention.output.LayerNorm.weight requires\_grad= False











```

bert.encoder.layer.22.attention.output.LayerNorm.bias requires_grad= False
bert.encoder.layer.22.intermediate.dense.weight requires_grad= False
bert.encoder.layer.22.intermediate.dense.bias requires_grad= False
bert.encoder.layer.22.output.dense.weight requires_grad= False
bert.encoder.layer.22.output.dense.bias requires_grad= False
bert.encoder.layer.22.output.LayerNorm.weight requires_grad= False
bert.encoder.layer.22.output.LayerNorm.bias requires_grad= False
bert.encoder.layer.23.attention.self.query.weight requires_grad= True
bert.encoder.layer.23.attention.self.query.bias requires_grad= True
bert.encoder.layer.23.attention.self.key.weight requires_grad= True
bert.encoder.layer.23.attention.self.key.bias requires_grad= True
bert.encoder.layer.23.attention.self.value.weight requires_grad= True
bert.encoder.layer.23.attention.self.value.bias requires_grad= True
bert.encoder.layer.23.attention.output.dense.weight requires_grad= True
bert.encoder.layer.23.attention.output.dense.bias requires_grad= True
bert.encoder.layer.23.attention.output.LayerNorm.weight requires_grad= True
bert.encoder.layer.23.attention.output.LayerNorm.bias requires_grad= True
bert.encoder.layer.23.intermediate.dense.weight requires_grad= True
bert.encoder.layer.23.intermediate.dense.bias requires_grad= True
bert.encoder.layer.23.output.dense.weight requires_grad= True
bert.encoder.layer.23.output.dense.bias requires_grad= True
bert.encoder.layer.23.output.LayerNorm.weight requires_grad= True
bert.encoder.layer.23.output.LayerNorm.bias requires_grad= True
bert.pooler.dense.weight requires_grad= True
bert.pooler.dense.bias requires_grad= True
classifier.weight requires_grad= True
classifier.bias requires_grad= True

```

```
[ ]: model.resize_token_embeddings(len(tokenizer))
```

```
[ ]: Embedding(28996, 1024, padding_idx=0)
```

```

[ ]: # Train & Evaluate
trained_model, trainer_obj = train_transformer_model(
    model = model,
    tokenizer = tokenizer,
    train_dataset = train_data_hf,
    val_dataset = val_data_hf,
    output_dir = dir_results,
    num_epochs = num_epochs,
    batch_size = size_batch,
    lr = learning_rate,
    weight_decay = regularization_weight_decay)
metrics = trainer_obj.evaluate()
print("Validation metrics:", metrics)
test_metrics = trainer_obj.evaluate(test_data_hf) if test_data_hf else None
print("Test metrics:", test_metrics)

```

```

/usr/local/lib/python3.11/dist-packages/transformers/training_args.py:1611:
FutureWarning: `evaluation_strategy` is deprecated and will be removed in
version 4.46 of Transformers. Use `eval_strategy` instead
    warnings.warn(
<ipython-input-20-81222a87c90b>:31: FutureWarning: `tokenizer` is deprecated and
will be removed in version 5.0.0 for `Trainer.__init__`. Use `processing_class`
instead.
    trainer = Trainer(

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Validation metrics: {'eval_loss': 0.7082364559173584, 'eval_accuracy': 0.572,
'eval_precision': 0.5106382978723404, 'eval_recall': 0.2222222222222222,
'eval_f1': 0.3096774193548387, 'eval_runtime': 2.2694,
'eval_samples_per_second': 110.16, 'eval_steps_per_second': 0.881, 'epoch': 1.0}
Test metrics: {'eval_loss': 0.7516778707504272, 'eval_accuracy': 0.5,
'eval_precision': 0.5217391304347826, 'eval_recall': 0.1889763779527559,
'eval_f1': 0.2774566473988439, 'eval_runtime': 2.4586,
'eval_samples_per_second': 101.685, 'eval_steps_per_second': 0.813, 'epoch':
1.0}

```

```

[ ]: # save model checkpoint
# timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
pacific_time = datetime.now(zoneinfo.ZoneInfo("America/Los_Angeles"))
timestamp = pacific_time.isoformat()
model_save_path = os.path.join(dir_models,
    f"{x_task}_{named_model}_{y_col}_{timestamp}")
trainer_obj.save_model(model_save_path)
print(f"Model checkpoint saved to: {model_save_path}")
# log experiment results
experiment_info = {
    "model_name": named_model,
    "learning_rate": learning_rate,
    "epochs": num_epochs,
    "batch_size": size_batch,
    "weight_decay": regularization_weight_decay,
    "x_task": x_task,
    "x_col": x_col,
    "y_col": y_col,
    "layers_to_unfreeze": layers_to_unfreeze}
model_info = gather_model_details(trained_model)
all_run_metrics = gather_all_run_metrics(
    trainer=trainer_obj,
    train_dataset=train_data_hf,
    val_dataset=val_data_hf,
    test_dataset=test_data_hf)
log_experiment_results_json(

```

```

    experiment_meta=experiment_info,
    model_details=model_info,
    run_metrics=all_run_metrics,
    log_file=log_filepath)
print(f"EXPERIMENT LOGGED TO: {log_filepath}")

```

Model checkpoint saved to: /content/drive/MyDrive/266-final/models/multi\_bert-large-cased\_binary\_complexity\_2025-04-13T18:54:16.408669-07:00

<IPython.core.display.HTML object>

EXPERIMENT LOGGED TO:  
/content/drive/MyDrive/266-final/results/experiment\_runs.txt

```

[ ]: prediction_output = trainer_obj.predict(test_data_hf)
raw_predictions = prediction_output.predictions
true_labels = prediction_output.label_ids
preds = np.argmax(raw_predictions, axis=1)
mismatch_indices = np.where(preds != true_labels)[0]

error_rows = []
for idx in mismatch_indices:
    text_value = df_test.iloc[idx][x_col]
    true_label_val = true_labels[idx]
    pred_label_val = preds[idx]

    error_rows.append({
        "hf_index": idx,
        "text": text_value,
        "true_label": true_label_val,
        "predicted_label": pred_label_val
    })

error_df = pd.DataFrame(error_rows)
df_test_for_merge = df_test.copy()
df_test_for_merge["error_matching_prefix"] = df_test_for_merge[x_col].str[:50]
# df_test_for_merge.drop(columns=[x_col], inplace=True)

error_df["error_matching_prefix"] = error_df["text"].str[:50]
error_df = error_df.merge(
    df_test_for_merge,
    on="error_matching_prefix",
    how="left",
    suffixes=("", "_source"))

error_df.to_csv("bert-base-cased_mismatches.csv", index=False)

# print("Number of misclassified samples:", len(error_df))

```

```

print("\nMerged error_df with extra columns:")
# display(error_df.head(15))

# print("\nConfusion Matrix:")
# cm = confusion_matrix(true_labels, preds)
# plt.figure(figsize=(8, 6))
# sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=True,
#             xticklabels=["Predicted 0", "Predicted 1"],
#             yticklabels=["True 0", "True 1"])
# plt.xlabel('Predicted Label')
# plt.ylabel('True Label')
# plt.title('Confusion Matrix')
# plt.tight_layout()
# plt.show()
# print("confusion matrix metrics: \n", cm)
# error_save_path = os.path.join(dir_results,
#     ↪ f"{x_task}_{named_model}_{y_col}_{x_col}_errors.csv")
# error_df.to_csv(error_save_path, index=False)
# print("Result saved to results directory.")
prediction_output = trainer_obj.predict(test_data_hf)
raw_predictions = prediction_output.predictions
true_labels = prediction_output.label_ids
preds = np.argmax(raw_predictions, axis=1)
df_test["true_label"] = true_labels
df_test["predicted_label"] = preds
df_test["is_incorrect"] = (df_test["predicted_label"] != df_test["true_label"])
df_test["avg_embedding"] = None
device = next(model.parameters()).device
for i in range(len(df_test)):
    text_value = df_test.iloc[i][x_col]
    e = tokenizer(text_value, return_tensors="pt", truncation=True,
    ↪ max_length=512).to(device)
    with torch.no_grad():
        emb = model.bert.embeddings.word_embeddings(e["input_ids"]).mean(dim=1).
        ↪ squeeze().cpu().numpy()
    df_test.at[i, "avg_embedding"] = emb
cm = confusion_matrix(df_test["true_label"], df_test["predicted_label"])
plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=True,
    ↪ xticklabels=["Predicted 0", "Predicted 1"], yticklabels=["True 0", "True 1"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.tight_layout()
plt.show()
# print("confusion matrix metrics:\n", cm)

```

```

df_plot = df_test.dropna(subset=["avg_embedding"]).copy()
embeddings = np.stack(df_plot["avg_embedding"].values)
pca = PCA(n_components=3)
reduced = pca.fit_transform(embeddings)
df_plot["pca_x"] = reduced[:,0]
df_plot["pca_y"] = reduced[:,1]
df_plot["pca_z"] = reduced[:,2]
colors = {
    ("bible", True): "red",
    ("bible", False): "orange",
    ("europarl", True): "yellow",
    ("europarl", False): "green",
    ("biomed", True): "blue",
    ("biomed", False): "purple"}
fig = plt.figure(figsize=(13,13))
ax = fig.add_subplot(111, projection='3d')
fig.set_facecolor("white")
ax.set_facecolor("white")
ax.xaxis._axinfo["grid"]["color"] = "gray"
ax.yaxis._axinfo["grid"]["color"] = "gray"
ax.zaxis._axinfo["grid"]["color"] = "gray"
for (corp, incorr), color in colors.items():
    subset = df_plot[(df_plot["corpus"]==corp) &
    (df_plot["is_incorrect"]==incorr)]
    ax.scatter(subset["pca_x"], subset["pca_y"], subset["pca_z"], c=color,
    s=20, alpha=0.8, label=f"{corp}, incorrect={incorr}")
ax.set_title("Average Embedding Value of Predictions, by Corpus", color="black")
ax.set_xlabel("PC1", color="black")
ax.set_ylabel("PC2", color="black")
ax.set_zlabel("PC3", color="black")
ax.legend(loc="best")
plt.show()
if "corpus" in df_test.columns:
    freqs = df_test["corpus"].value_counts()
    print("\nFrequency counts of corpus:", freqs)
    err_df = df_test[df_test["is_incorrect"]==True]
    corr_df = df_test[df_test["is_incorrect"]==False]
    err_counts = err_df["corpus"].value_counts()
    corr_counts = corr_df["corpus"].value_counts()
    print("\nCounts of corpus in misclassified:", err_counts)
    print("\nCounts of corpus in correctly classified:", corr_counts)
    print("\nProportions of corpus in misclassified:", err_counts/err_counts.
    sum())
    print("\nProportions of corpus in correctly classified:", corr_counts/
    corr_counts.sum())
    grouped_all = df_test.groupby("corpus")["avg_embedding"].apply(lambda x: np.
    mean(np.stack(x.values), axis=0))

```

```

    grouped_err = err_df.groupby("corpus")["avg_embedding"].apply(lambda x: np.
↳mean(np.stack(x.values), axis=0))
    grouped_corr = corr_df.groupby("corpus")["avg_embedding"].apply(lambda x:
↳np.mean(np.stack(x.values), axis=0))
    # print("\nAvg embedding of each subcorpus overall:", grouped_all)
    # print("\nAvg embedding of each subcorpus misclassified:", grouped_err)
    # print("\nAvg embedding of each subcorpus correctly classified:",
↳grouped_corr)
err_stack = np.stack(df_test[df_test["is_incorrect"]==True]["avg_embedding"].
↳values)
corr_stack = np.stack(df_test[df_test["is_incorrect"]==False]["avg_embedding"].
↳values)
# print("\nAvg embedding of records predicted incorrectly:", err_stack.
↳mean(axis=0))
# print("Avg embedding of records predicted correctly:", corr_stack.
↳mean(axis=0))
error_df = df_test[df_test["is_incorrect"]==True].copy()
error_df.to_csv("misclassified_with_all_columns.csv", index=False)
error_save_path = os.path.join(dir_results,
↳f"{x_task}_{named_model}_{y_col}_{x_col}_errors.csv")
df_test.to_csv(error_save_path, index=False)
for corp in df_test["corpus"].unique():
    subset = df_test[df_test["corpus"]==corp]
    emb_true = subset[subset["is_incorrect"]==False]["avg_embedding"]
    emb_false = subset[subset["is_incorrect"]==True]["avg_embedding"]
    if len(emb_true)==0 or len(emb_false)==0:
        print(f"No valid data for subcorpus '{corp}'")
        continue
    p = np.mean(np.stack(emb_true.values), axis=0)
    q = np.mean(np.stack(emb_false.values), axis=0)
    p_exp = np.exp(p - np.max(p))
    q_exp = np.exp(q - np.max(q))
    p_sum = p_exp.sum()
    q_sum = q_exp.sum()
    if p_sum<=0 or q_sum<=0:
        print(f"Cannot form valid distributions for subcorpus '{corp}'")
        continue
    p_dist = p_exp / p_sum
    q_dist = q_exp / q_sum
    kl_pq = entropy(p_dist, q_dist)
    kl_qp = entropy(q_dist, p_dist)
    kl_sym = 0.5*(kl_pq + kl_qp)
    print(f"Subcorpus '{corp}' symmetric KL divergence: {kl_sym}")
error_save_path = os.path.join(dir_results,
↳f"{x_task}_{named_model}_{y_col}_{x_col}_errors.csv")
df_test.to_csv(error_save_path, index=False)

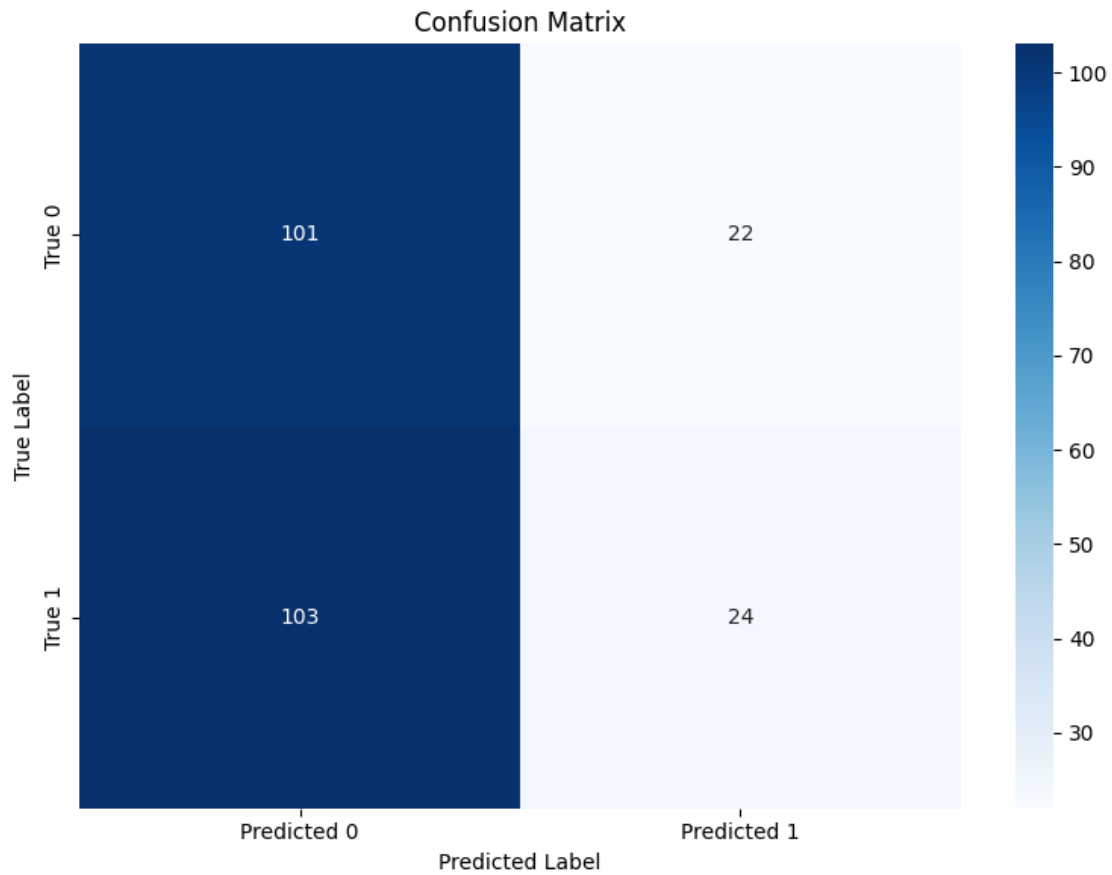
```


```
print("model results saved")
```

<IPython.core.display.HTML object>

Merged error\_df with extra columns:

<IPython.core.display.HTML object>





```
Frequency counts of corpus: corpus
```

```
bible      88
```

```
europarl   87
```

```
biomed     75
```

```
Name: count, dtype: int64
```

```
Counts of corpus in misclassified: corpus
```

```
europarl   43
```

```
biomed     42
```

```
bible      40
```

```
Name: count, dtype: int64
```



Counts of corpus in correctly classified: corpus

bible 48

europarl 44

biomed 33

Name: count, dtype: int64

Proportions of corpus in misclassified: corpus

europarl 0.344

biomed 0.336

bible 0.320

Name: count, dtype: float64

Proportions of corpus in correctly classified: corpus

bible 0.384

europarl 0.352

biomed 0.264

Name: count, dtype: float64

Subcorpus 'bible' symmetric KL divergence: 1.888256876176576e-06

Subcorpus 'biomed' symmetric KL divergence: 1.6649774807960557e-06

Subcorpus 'europarl' symmetric KL divergence: 3.837849675757952e-06

model results saved

## Result

### 0.2.5 answerdotai/ModernBERT-base

```
[ ]: # Define Experiment Parameters
# named_model = "bert-base-cased"
# named_model = "roberta-base"
# named_model = "bert-large-cased"
# named_model = "roberta-large"
named_model = "answerdotai/ModernBERT-base" # modern bert
#####
regularization_weight_decay = 0.5
learning_rate = 5e-6
size_batch = 128
length_max = 128
num_epochs = 1
#####
# x_col = "sentence"
x_col = "sentence_no_contractions"
# x_col = "pos_sequence"
# x_col = "dep_sequence"
# x_col = "morph_sequence"
# x_col = "snc_pos_seq"
# x_col = "snc_pos_alt"
# x_col = "snc_morph_seq"
```

```

# x_col = "snc_morph_alt"
# x_col = "snc_dep_seq"
# x_col = "snc_dep_alt"
# x_col = "snc_morph_complexity_value"
#####
y_col = "binary_complexity"
# y_col = "complexity"
#####
# x_task = "single"
x_task = "multi"
if x_task == "single":
    df_train = train_single_df
    df_val = trial_val_single_df
    df_test = test_single_df
else:
    df_train = train_multi_df
    df_val = trial_val_multi_df
    df_test = test_multi_df
#####
# Tokenize & Prepare Datasets
train_data_hf = prepare_dataset(
    df_train,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
val_data_hf = prepare_dataset(
    df_val,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
test_data_hf = prepare_dataset(
    df_test,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
print("Datasets prepared. Sample from train_data_hf:\n", train_data_hf[10])
# print("Datasets prepared. Sample from train_data_hf:\n", val_data_hf[10])
# print("Datasets prepared. Sample from train_data_hf:\n", test_data_hf[10])
#####
# custom_config = BertConfig.from_pretrained("roberta-base")
# custom_config.hidden_act = "gelu" # alts: "relu" "silu"
# custom_config.attention_probs_dropout_prob = 0.1
# custom_config.hidden_dropout_prob = 0.1
# custom_config.gradient_checkpointing = False

```



Some weights of ModernBertForSequenceClassification were not initialized from the model checkpoint at answerdotai/ModernBERT-base and are newly initialized: ['classifier.bias', 'classifier.weight']  
 You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
=====
answerdotai/ModernBERT-base :
=====
num_parameters: 149606402
num_trainable_parameters at load: 149606402
=====
model lineage: {'type': 'huggingface_hub', 'path': 'answerdotai/ModernBERT-base', 'timestamp': '2025-04-14 01:54:41'}
=====
```

```
[ ]: print(model)
```

```
ModernBertForSequenceClassification(
  (model): ModernBertModel(
    (embeddings): ModernBertEmbeddings(
      (tok_embeddings): Embedding(50368, 768, padding_idx=50283)
      (norm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (drop): Dropout(p=0.0, inplace=False)
    )
    (layers): ModuleList(
      (0): ModernBertEncoderLayer(
        (attn_norm): Identity()
        (attn): ModernBertAttention(
          (Wqkv): Linear(in_features=768, out_features=2304, bias=False)
          (rotary_emb): ModernBertRotaryEmbedding()
          (Wo): Linear(in_features=768, out_features=768, bias=False)
          (out_drop): Identity()
        )
        (mlp_norm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
        (mlp): ModernBertMLP(
          (Wi): Linear(in_features=768, out_features=2304, bias=False)
          (act): GELUActivation()
          (drop): Dropout(p=0.0, inplace=False)
          (Wo): Linear(in_features=1152, out_features=768, bias=False)
        )
      )
    )
    (1-21): 21 x ModernBertEncoderLayer(
      (attn_norm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
      (attn): ModernBertAttention(
        (Wqkv): Linear(in_features=768, out_features=2304, bias=False)
        (rotary_emb): ModernBertRotaryEmbedding()
        (Wo): Linear(in_features=768, out_features=768, bias=False)
        (out_drop): Identity()
      )
    )
  )
)
```

```

    )
    (mlp_norm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
    (mlp): ModernBertMLP(
      (Wi): Linear(in_features=768, out_features=2304, bias=False)
      (act): GELUActivation()
      (drop): Dropout(p=0.0, inplace=False)
      (Wo): Linear(in_features=1152, out_features=768, bias=False)
    )
  )
  (final_norm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
)
(head): ModernBertPredictionHead(
  (dense): Linear(in_features=768, out_features=768, bias=False)
  (act): GELUActivation()
  (norm): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
)
(drop): Dropout(p=0.0, inplace=False)
(classifier): Linear(in_features=768, out_features=2, bias=True)
)

```

```

[ ]: # for name, param in model.named_parameters():
#     print(name, "requires_grad=", param.requires_grad)

```

```

[ ]: # # Inspect the attention_mask tensor for the first few samples
# for i in range(5):
#     print(train_data_hf[i]['attention_mask'])

```

```

[ ]: #####
layers_to_unfreeze = [
    "model.layers.21.attn_norm.weight",
    "model.layers.21.attn.Wqkv.weight",
    "model.layers.21.attn.Wo.weight",
    "model.layers.21.mlp_norm.weight",
    "model.layers.21.mlp.Wi.weight",
    "model.layers.21.mlp.Wo.weight",
    "model.final_norm.weight",
    "head.dense.weight",
    "head.norm.weight",
    "classifier.weight",
    "classifier.bias"]
freeze_unfreeze_layers(model, layers_to_unfreeze=layers_to_unfreeze)
print(model.config)
print("=====")
print("num_parameters:", model.num_parameters())
print("num_trainable_parameters:", model.num_parameters(only_trainable=True))
print("=====")

```

```
#####
print("Experiment configuration used with this experiment:")
print("model used:", named_model)
print("learning rate used:", learning_rate)
print("number of epochs:", num_epochs)
print("maximum sequence length:", length_max)
print("batch size used:", size_batch)
print("regularization value:", regularization_weight_decay)
print("outcome variable:", y_col)
print("task:", x_task)
print("input column:", x_col)
#####
print("=====")
print("num_trainable_parameters:", model.num_parameters(only_trainable=True))
```

```
ModernBertConfig {
  "_attn_implementation_autoset": true,
  "architectures": [
    "ModernBertForMaskedLM"
  ],
  "attention_bias": false,
  "attention_dropout": 0.0,
  "bos_token_id": 50281,
  "classifier_activation": "gelu",
  "classifier_bias": false,
  "classifier_dropout": 0.0,
  "classifier_pooling": "mean",
  "cls_token_id": 50281,
  "decoder_bias": true,
  "deterministic_flash_attn": false,
  "embedding_dropout": 0.0,
  "eos_token_id": 50282,
  "global_attn_every_n_layers": 3,
  "global_rope_theta": 160000.0,
  "gradient_checkpointing": false,
  "hidden_activation": "gelu",
  "hidden_size": 768,
  "initializer_cutoff_factor": 2.0,
  "initializer_range": 0.02,
  "intermediate_size": 1152,
  "layer_norm_eps": 1e-05,
  "local_attention": 128,
  "local_rope_theta": 10000.0,
  "max_position_embeddings": 8192,
  "mlp_bias": false,
  "mlp_dropout": 0.0,
  "model_type": "modernbert",
  "norm_bias": false,
```

```

    "norm_eps": 1e-05,
    "num_attention_heads": 12,
    "num_hidden_layers": 22,
    "pad_token_id": 50283,
    "position_embedding_type": "absolute",
    "reference_compile": null,
    "repad_logits_with_grad": false,
    "sep_token_id": 50282,
    "sparse_pred_ignore_index": -100,
    "sparse_prediction": false,
    "torch_dtype": "float32",
    "transformers_version": "4.50.3",
    "vocab_size": 50368
}

```

```

=====

```

```

num_parameters: 149606402
num_trainable_parameters: 5607938

```

```

=====

```

Experiment configuration used with this experiment:

model used: answerdotai/ModernBERT-base

learning rate used: 5e-06

number of epochs: 1

maximum sequence length: 128

batch size used: 128

regularization value: 0.5

outcome variable: binary\_complexity

task: multi

input column: sentence\_no\_contractions

```

=====

```

```

num_trainable_parameters: 5607938

```

```

[ ]: # for name, param in model.named_parameters():
    #     print(name, "requires_grad=", param.requires_grad)

```

```

[ ]: # model.resize_token_embeddings(len(tokenizer))

```

```

[ ]: # Train & Evaluate
trained_model, trainer_obj = train_transformer_model(
    model = model,
    tokenizer = tokenizer,
    train_dataset = train_data_hf,
    val_dataset = val_data_hf,
    output_dir = dir_results,
    num_epochs = num_epochs,
    batch_size = size_batch,
    lr = learning_rate,

```

```

weight_decay = regularization_weight_decay)
metrics = trainer_obj.evaluate()
print("Validation metrics:", metrics)
test_metrics = trainer_obj.evaluate(test_data_hf) if test_data_hf else None
print("Test metrics:", test_metrics)

```

```

/usr/local/lib/python3.11/dist-packages/transformers/training_args.py:1611:
FutureWarning: `evaluation_strategy` is deprecated and will be removed in
version 4.46 of Transformers. Use `eval_strategy` instead
warnings.warn(
<ipython-input-20-81222a87c90b>:31: FutureWarning: `tokenizer` is deprecated and
will be removed in version 5.0.0 for `Trainer.__init__`. Use `processing_class`
instead.
trainer = Trainer(
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
Validation metrics: {'eval_loss': 0.9965472221374512, 'eval_accuracy': 0.448,
'eval_precision': 0.42718446601941745, 'eval_recall': 0.8148148148148148,
'eval_f1': 0.5605095541401274, 'eval_runtime': 1.589, 'eval_samples_per_second':
157.33, 'eval_steps_per_second': 1.259, 'epoch': 1.0}
Test metrics: {'eval_loss': 0.9817145466804504, 'eval_accuracy': 0.504,
'eval_precision': 0.507537688442211, 'eval_recall': 0.7952755905511811,
'eval_f1': 0.6196319018404908, 'eval_runtime': 1.6323,
'eval_samples_per_second': 153.154, 'eval_steps_per_second': 1.225, 'epoch':
1.0}

```

```

[ ]: # save model checkpoint
# timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
pacific_time = datetime.now(zoneinfo.ZoneInfo("America/Los_Angeles"))
timestamp = pacific_time.isoformat()
model_save_path = os.path.join(dir_models,
    ↪f"{x_task}_{named_model}_{y_col}_{timestamp}")
trainer_obj.save_model(model_save_path)
print(f"Model checkpoint saved to: {model_save_path}")
# log experiment results
experiment_info = {
    "model_name": named_model,
    "learning_rate": learning_rate,
    "epochs": num_epochs,
    "batch_size": size_batch,
    "weight_decay": regularization_weight_decay,
    "x_task": x_task,
    "x_col": x_col,
    "y_col": y_col,
    "layers_to_unfreeze": layers_to_unfreeze}
model_info = gather_model_details(trained_model)

```



```

all_run_metrics = gather_all_run_metrics(
    trainer=trainer_obj,
    train_dataset=train_data_hf,
    val_dataset=val_data_hf,
    test_dataset=test_data_hf)
log_experiment_results_json(
    experiment_meta=experiment_info,
    model_details=model_info,
    run_metrics=all_run_metrics,
    log_file=log_filepath)
print(f"EXPERIMENT LOGGED TO: {log_filepath}")

```

Model checkpoint saved to:

/content/drive/MyDrive/266-final/models/multi\_answerdotai/ModernBERT-  
base\_binary\_complexity\_2025-04-13T18:54:57.678930-07:00

<IPython.core.display.HTML object>

EXPERIMENT LOGGED TO:

/content/drive/MyDrive/266-final/results/experiment\_runs.txt

```

[ ]: prediction_output = trainer_obj.predict(test_data_hf)
raw_predictions = prediction_output.predictions
true_labels = prediction_output.label_ids
preds = np.argmax(raw_predictions, axis=1)
mismatch_indices = np.where(preds != true_labels)[0]

error_rows = []
for idx in mismatch_indices:
    text_value = df_test.iloc[idx][x_col]
    true_label_val = true_labels[idx]
    pred_label_val = preds[idx]

    error_rows.append({
        "hf_index": idx,
        "text": text_value,
        "true_label": true_label_val,
        "predicted_label": pred_label_val
    })

error_df = pd.DataFrame(error_rows)
df_test_for_merge = df_test.copy()
df_test_for_merge["error_matching_prefix"] = df_test_for_merge[x_col].str[:50]
# df_test_for_merge.drop(columns=[x_col], inplace=True)

error_df["error_matching_prefix"] = error_df["text"].str[:50]
error_df = error_df.merge(
    df_test_for_merge,

```

```

        on="error_matching_prefix",
        how="left",
        suffixes=("", "_source"))

error_df.to_csv("bert-base-cased_mismatches.csv", index=False)

# print("Number of misclassified samples:", len(error_df))
print("\nMerged error_df with extra columns:")
# display(error_df.head(15))

# print("\nConfusion Matrix:")
# cm = confusion_matrix(true_labels, preds)
# plt.figure(figsize=(8, 6))
# sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=True,
#             xticklabels=["Predicted 0", "Predicted 1"],
#             yticklabels=["True 0", "True 1"])
# plt.xlabel('Predicted Label')
# plt.ylabel('True Label')
# plt.title('Confusion Matrix')
# plt.tight_layout()
# plt.show()
# print("confusion matrix metrics: \n", cm)
# error_save_path = os.path.join(dir_results,
#     ↳f"{x_task}_{named_model}_{y_col}_{x_col}_errors.csv")
# error_df.to_csv(error_save_path, index=False)
# print("Result saved to results directory.")
prediction_output = trainer_obj.predict(test_data_hf)
raw_predictions = prediction_output.predictions
true_labels = prediction_output.label_ids
preds = np.argmax(raw_predictions, axis=1)
df_test["true_label"] = true_labels
df_test["predicted_label"] = preds
df_test["is_incorrect"] = (df_test["predicted_label"] != df_test["true_label"])
df_test["avg_embedding"] = None
device = next(model.parameters()).device
for i in range(len(df_test)):
    text_value = df_test.iloc[i][x_col]
    e = tokenizer(text_value, return_tensors="pt", truncation=True,
        ↳max_length=512).to(device)
    with torch.no_grad():
        # emb = model.bert.embeddings.word_embeddings(e["input_ids"]).
        ↳mean(dim=1).squeeze().cpu().numpy()
        emb = model.model.embeddings.tok_embeddings(e["input_ids"]).mean(dim=1).
        ↳squeeze().cpu().numpy()
    df_test.at[i, "avg_embedding"] = emb
cm = confusion_matrix(df_test["true_label"], df_test["predicted_label"])
plt.figure(figsize=(8,6))

```

```

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=True,
    xticklabels=["Predicted 0", "Predicted 1"], yticklabels=["True 0", "True 1"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.tight_layout()
plt.show()
# print("confusion matrix metrics:\n", cm)

df_plot = df_test.dropna(subset=["avg_embedding"]).copy()
embeddings = np.stack(df_plot["avg_embedding"].values)
pca = PCA(n_components=3)
reduced = pca.fit_transform(embeddings)
df_plot["pca_x"] = reduced[:,0]
df_plot["pca_y"] = reduced[:,1]
df_plot["pca_z"] = reduced[:,2]
colors = {
    ("bible", True): "red",
    ("bible", False): "orange",
    ("europarl", True): "yellow",
    ("europarl", False): "green",
    ("biomed", True): "blue",
    ("biomed", False): "purple"}
fig = plt.figure(figsize=(13,13))
ax = fig.add_subplot(111, projection='3d')
fig.set_facecolor("white")
ax.set_facecolor("white")
ax.xaxis._axinfo["grid"]["color"] = "gray"
ax.yaxis._axinfo["grid"]["color"] = "gray"
ax.zaxis._axinfo["grid"]["color"] = "gray"
for (corp, incorr), color in colors.items():
    subset = df_plot[(df_plot["corpus"]==corp) &
    (df_plot["is_incorrect"]==incorr)]
    ax.scatter(subset["pca_x"], subset["pca_y"], subset["pca_z"], c=color,
    s=20, alpha=0.8, label=f"{corp}, incorrect={incorr}")
ax.set_title("Average Embedding Value of Predictions, by Corpus", color="black")
ax.set_xlabel("PC1", color="black")
ax.set_ylabel("PC2", color="black")
ax.set_zlabel("PC3", color="black")
ax.legend(loc="best")
plt.show()
if "corpus" in df_test.columns:
    freqs = df_test["corpus"].value_counts()
    print("\nFrequency counts of corpus:", freqs)
    err_df = df_test[df_test["is_incorrect"]==True]
    corr_df = df_test[df_test["is_incorrect"]==False]
    err_counts = err_df["corpus"].value_counts()

```

```

corr_counts = corr_df["corpus"].value_counts()
print("\nCounts of corpus in misclassified:", err_counts)
print("\nCounts of corpus in correctly classified:", corr_counts)
print("\nProportions of corpus in misclassified:", err_counts/err_counts.
↳sum())
    print("\nProportions of corpus in correctly classified:", corr_counts/
↳corr_counts.sum())
    grouped_all = df_test.groupby("corpus")["avg_embedding"].apply(lambda x: np.
↳mean(np.stack(x.values), axis=0))
    grouped_err = err_df.groupby("corpus")["avg_embedding"].apply(lambda x: np.
↳mean(np.stack(x.values), axis=0))
    grouped_corr = corr_df.groupby("corpus")["avg_embedding"].apply(lambda x:
↳np.mean(np.stack(x.values), axis=0))
    # print("\nAvg embedding of each subcorpus overall:", grouped_all)
    # print("\nAvg embedding of each subcorpus misclassified:", grouped_err)
    # print("\nAvg embedding of each subcorpus correctly classified:",
↳grouped_corr)
err_stack = np.stack(df_test[df_test["is_incorrect"]==True]["avg_embedding"].
↳values)
corr_stack = np.stack(df_test[df_test["is_incorrect"]==False]["avg_embedding"].
↳values)
# print("\nAvg embedding of records predicted incorrectly:", err_stack.
↳mean(axis=0))
# print("Avg embedding of records predicted correctly:", corr_stack.
↳mean(axis=0))
error_df = df_test[df_test["is_incorrect"]==True].copy()
# error_df.to_csv("misclassified_with_all_columns.csv", index=False)
# error_save_path = os.path.join(dir_results,
↳f"{x_task}_modernbert-base_{y_col}_{x_col}_errors.csv")
df_test.to_csv(error_save_path, index=False)
for corp in df_test["corpus"].unique():
    subset = df_test[df_test["corpus"]==corp]
    emb_true = subset[subset["is_incorrect"]==False]["avg_embedding"]
    emb_false = subset[subset["is_incorrect"]==True]["avg_embedding"]
    if len(emb_true)==0 or len(emb_false)==0:
        print(f"No valid data for subcorpus '{corp}'")
        continue
    p = np.mean(np.stack(emb_true.values), axis=0)
    q = np.mean(np.stack(emb_false.values), axis=0)
    p_exp = np.exp(p - np.max(p))
    q_exp = np.exp(q - np.max(q))
    p_sum = p_exp.sum()
    q_sum = q_exp.sum()
    if p_sum<=0 or q_sum<=0:
        print(f"Cannot form valid distributions for subcorpus '{corp}'")
        continue

```

```

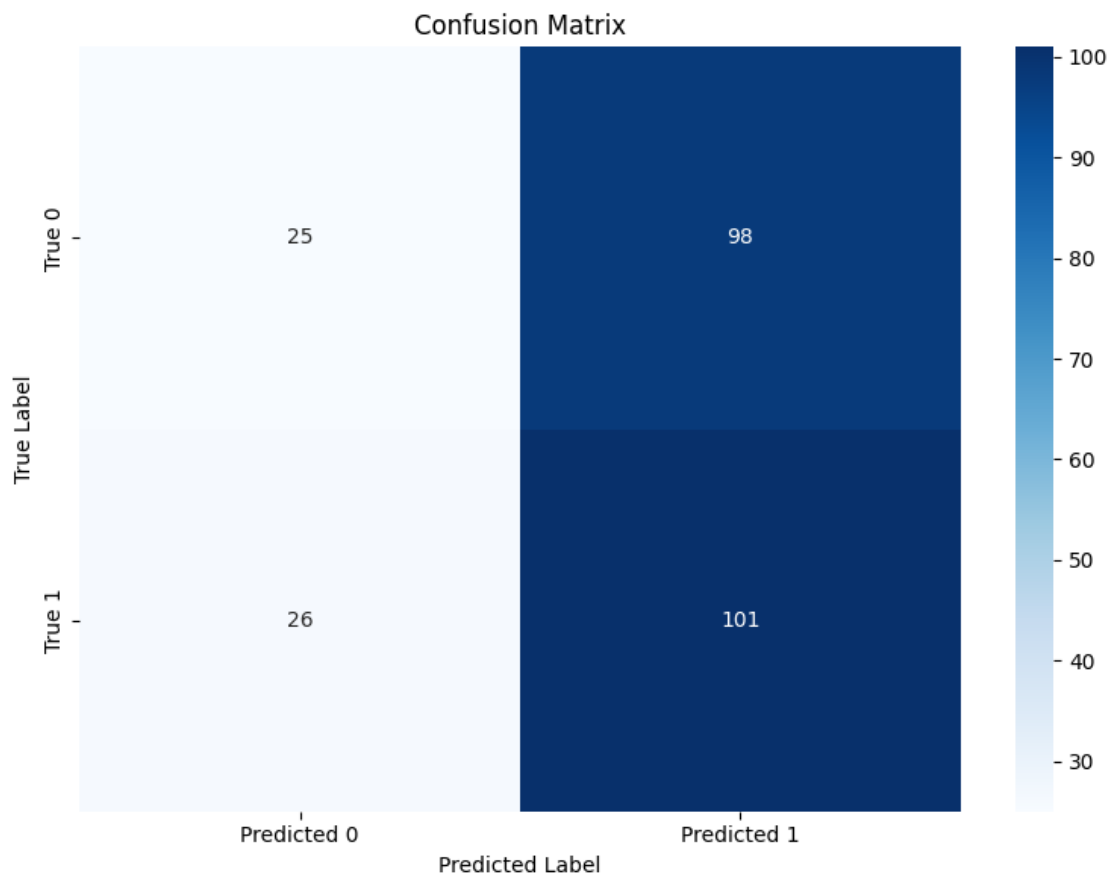
p_dist = p_exp / p_sum
q_dist = q_exp / q_sum
kl_pq = entropy(p_dist, q_dist)
kl_qp = entropy(q_dist, p_dist)
kl_sym = 0.5*(kl_pq + kl_qp)
print(f"Subcorpus '{corp}' symmetric KL divergence: {kl_sym}")
error_save_path = os.path.join(dir_results,
    f"{x_task}_modernbert-base_{y_col}_{x_col}_errors.csv")
df_test.to_csv(error_save_path, index=False)
print("model results saved")

```


<IPython.core.display.HTML object>

Merged error\_df with extra columns:

<IPython.core.display.HTML object>



Average Embedding Value of Predictions, by Corpus



Frequency counts of corpus: corpus

bible 88

europarl 87

biomed 75

Name: count, dtype: int64

Counts of corpus in misclassified: corpus

europarl 48

bible 45

biomed 31

Name: count, dtype: int64

Counts of corpus in correctly classified: corpus

biomed 44

bible 43

europarl 39

Name: count, dtype: int64

Proportions of corpus in misclassified: corpus

europarl 0.387097

bible 0.362903

biomed 0.250000

Name: count, dtype: float64

Proportions of corpus in correctly classified: corpus

biomed 0.349206

bible 0.341270

europarl 0.309524

Name: count, dtype: float64

Subcorpus 'bible' symmetric KL divergence: 8.12416918642706e-06

Subcorpus 'biomed' symmetric KL divergence: 6.612745022528127e-06

Subcorpus 'europarl' symmetric KL divergence: 1.2013919929424917e-05

model results saved

## Result

### 0.2.6 answerdotai/ModernBERT-large

```
[ ]: # Define Experiment Parameters
# named_model = "bert-base-cased"
# named_model = "roberta-base"
# named_model = "bert-large-cased"
# named_model = "roberta-large"
named_model = "answerdotai/ModernBERT-large" # modern bert
#####
regularization_weight_decay = 0.5
learning_rate = 5e-6
size_batch = 128
length_max = 128
num_epochs = 1
#####
# x_col = "sentence"
x_col = "sentence_no_contractions"
# x_col = "pos_sequence"
# x_col = "dep_sequence"
# x_col = "morph_sequence"
# x_col = "snc_pos_seq"
# x_col = "snc_pos_alt"
# x_col = "snc_morph_seq"
```

```

# x_col = "snc_morph_alt"
# x_col = "snc_dep_seq"
# x_col = "snc_dep_alt"
# x_col = "snc_morph_complexity_value"
#####
y_col = "binary_complexity"
# y_col = "complexity"
#####
# x_task = "single"
x_task = "multi"
if x_task == "single":
    df_train = train_single_df
    df_val = trial_val_single_df
    df_test = test_single_df
else:
    df_train = train_multi_df
    df_val = trial_val_multi_df
    df_test = test_multi_df
#####
# Tokenize & Prepare Datasets
train_data_hf = prepare_dataset(
    df_train,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
val_data_hf = prepare_dataset(
    df_val,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
test_data_hf = prepare_dataset(
    df_test,
    tokenizer,
    text_col=x_col,
    label_col=y_col,
    max_length=length_max)
print("Datasets prepared. Sample from train_data_hf:\n", train_data_hf[10])
# print("Datasets prepared. Sample from train_data_hf:\n", val_data_hf[10])
# print("Datasets prepared. Sample from train_data_hf:\n", test_data_hf[10])
#####
model, tokenizer = get_model_and_tokenizer(
    remote_model_name="answerdotai/ModernBERT-large",
    local_model_path=None,
    config=None)
#####

```





```

=====
answerdotai/ModernBERT-large :
=====
num_parameters: 395833346
num_trainable_parameters at load: 395833346
=====
model lineage: {'type': 'huggingface_hub', 'path': 'answerdotai/ModernBERT-
large', 'timestamp': '2025-04-14 01:55:15'}
=====

```

```
[ ]: print(model)
```

```

ModernBertForSequenceClassification(
  (model): ModernBertModel(
    (embeddings): ModernBertEmbeddings(
      (tok_embeddings): Embedding(50368, 1024, padding_idx=50283)
      (norm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
      (drop): Dropout(p=0.0, inplace=False)
    )
    (layers): ModuleList(
      (0): ModernBertEncoderLayer(
        (attn_norm): Identity()
        (attn): ModernBertAttention(
          (Wqkv): Linear(in_features=1024, out_features=3072, bias=False)
          (rotary_emb): ModernBertRotaryEmbedding()
          (Wo): Linear(in_features=1024, out_features=1024, bias=False)
          (out_drop): Identity()
        )
        (mlp_norm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
        (mlp): ModernBertMLP(
          (Wi): Linear(in_features=1024, out_features=5248, bias=False)
          (act): GELUActivation()
          (drop): Dropout(p=0.0, inplace=False)
          (Wo): Linear(in_features=2624, out_features=1024, bias=False)
        )
      )
      (1-27): 27 x ModernBertEncoderLayer(
        (attn_norm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
        (attn): ModernBertAttention(
          (Wqkv): Linear(in_features=1024, out_features=3072, bias=False)
          (rotary_emb): ModernBertRotaryEmbedding()
          (Wo): Linear(in_features=1024, out_features=1024, bias=False)
          (out_drop): Identity()
        )
        (mlp_norm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
        (mlp): ModernBertMLP(
          (Wi): Linear(in_features=1024, out_features=5248, bias=False)
          (act): GELUActivation()

```

```

        (drop): Dropout(p=0.0, inplace=False)
        (Wo): Linear(in_features=2624, out_features=1024, bias=False)
    )
)
)
(final_norm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
)
(head): ModernBertPredictionHead(
    (dense): Linear(in_features=1024, out_features=1024, bias=False)
    (act): GELUActivation()
    (norm): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
)
(drop): Dropout(p=0.0, inplace=False)
(classifier): Linear(in_features=1024, out_features=2, bias=True)
)

```

```

[ ]: # for name, param in model.named_parameters():
#     print(name, "requires_grad=", param.requires_grad)

```

```

[ ]: # # Inspect the attention_mask tensor for the first few samples
# for i in range(5):
#     print(train_data_hf[i]['attention_mask'])

```

```

[ ]: #####
layers_to_unfreeze = [
    "model.layers.27.attn_norm.weight",
    "model.layers.27.attn.Wqkv.weight",
    "model.layers.27.attn.Wo.weight",
    "model.layers.27.mlp_norm.weight",
    "model.layers.27.mlp.Wi.weight",
    "model.layers.27.mlp.Wo.weight",
    "model.final_norm.weight",
    "head.dense.weight",
    "head.norm.weight",
    "classifier.weight",
    "classifier.bias"
]
freeze_unfreeze_layers(model, layers_to_unfreeze=layers_to_unfreeze)
print(model.config)
print("=====")
print("num_parameters:", model.num_parameters())
print("num_trainable_parameters:", model.num_parameters(only_trainable=True))
print("=====")
#####
print("Experiment configuration used with this experiment:")
print("model used:", named_model)
print("learning rate used:", learning_rate)

```

```

print("number of epochs:", num_epochs)
print("maximum sequence length:", length_max)
print("batch size used:", size_batch)
print("regularization value:", regularization_weight_decay)
print("outcome variable:", y_col)
print("task:", x_task)
print("input column:", x_col)
#####
print("=====")
print("num_trainable_parameters:", model.num_parameters(only_trainable=True))

```

```

ModernBertConfig {
  "_attn_implementation_autoset": true,
  "architectures": [
    "ModernBertForMaskedLM"
  ],
  "attention_bias": false,
  "attention_dropout": 0.0,
  "bos_token_id": 50281,
  "classifier_activation": "gelu",
  "classifier_bias": false,
  "classifier_dropout": 0.0,
  "classifier_pooling": "mean",
  "cls_token_id": 50281,
  "decoder_bias": true,
  "deterministic_flash_attn": false,
  "embedding_dropout": 0.0,
  "eos_token_id": 50282,
  "global_attn_every_n_layers": 3,
  "global_rope_theta": 160000.0,
  "gradient_checkpointing": false,
  "hidden_activation": "gelu",
  "hidden_size": 1024,
  "initializer_cutoff_factor": 2.0,
  "initializer_range": 0.02,
  "intermediate_size": 2624,
  "layer_norm_eps": 1e-05,
  "local_attention": 128,
  "local_rope_theta": 10000.0,
  "max_position_embeddings": 8192,
  "mlp_bias": false,
  "mlp_dropout": 0.0,
  "model_type": "modernbert",
  "norm_bias": false,
  "norm_eps": 1e-05,
  "num_attention_heads": 16,
  "num_hidden_layers": 28,
  "pad_token_id": 50283,

```

```

    "position_embedding_type": "absolute",
    "reference_compile": null,
    "repad_logits_with_grad": false,
    "sep_token_id": 50282,
    "sparse_pred_ignore_index": -100,
    "sparse_prediction": false,
    "torch_dtype": "float32",
    "transformers_version": "4.50.3",
    "vocab_size": 50368
}

```

```

=====

```

```

num_parameters: 395833346
num_trainable_parameters: 13309954
=====

```

```

Experiment configuration used with this experiment:
model used: answerdotai/ModernBERT-large
learning rate used: 5e-06
number of epochs: 1
maximum sequence length: 128
batch size used: 128
regularization value: 0.5
outcome variable: binary_complexity
task: multi
input column: sentence_no_contractions
=====
num_trainable_parameters: 13309954

```

```

[ ]: # for name, param in model.named_parameters():
    #     print(name, "requires_grad=", param.requires_grad)

```

```

[ ]: # model.resize_token_embeddings(len(tokenizer))

```

```

[ ]: # Train & Evaluate
trained_model, trainer_obj = train_transformer_model(
    model = model,
    tokenizer = tokenizer,
    train_dataset = train_data_hf,
    val_dataset = val_data_hf,
    output_dir = dir_results,
    num_epochs = num_epochs,
    batch_size = size_batch,
    lr = learning_rate,
    weight_decay = regularization_weight_decay)
metrics = trainer_obj.evaluate()
print("Validation metrics:", metrics)
test_metrics = trainer_obj.evaluate(test_data_hf) if test_data_hf else None

```

```
print("Test metrics:", test_metrics)
```

```
/usr/local/lib/python3.11/dist-packages/transformers/training_args.py:1611:
FutureWarning: `evaluation_strategy` is deprecated and will be removed in
version 4.46 of Transformers. Use `eval_strategy` instead
  warnings.warn(
<ipython-input-20-81222a87c90b>:31: FutureWarning: `tokenizer` is deprecated and
will be removed in version 5.0.0 for `Trainer.__init__`. Use `processing_class`
instead.
  trainer = Trainer(
<IPython.core.display.HTML object>
W0414 01:55:25.351000 17035 torch/_dynamo/convert_frame.py:906] [1/8]
torch._dynamo hit config.cache_size_limit (8)
W0414 01:55:25.351000 17035 torch/_dynamo/convert_frame.py:906] [1/8]
function: 'compiled_mlp' (/usr/local/lib/python3.11/dist-
packages/transformers/models/modernbert/modeling_modernbert.py:552)
W0414 01:55:25.351000 17035 torch/_dynamo/convert_frame.py:906] [1/8] last
reason: 1/0: GLOBAL_STATE changed: grad_mode
W0414 01:55:25.351000 17035 torch/_dynamo/convert_frame.py:906] [1/8] To log all
recompilation reasons, use TORCH_LOGS="recompiles".
W0414 01:55:25.351000 17035 torch/_dynamo/convert_frame.py:906] [1/8] To
diagnose recompilation issues, see
https://pytorch.org/docs/main/torch.compiler\_troubleshooting.html.
<IPython.core.display.HTML object>
Validation metrics: {'eval_loss': 0.7263299822807312, 'eval_accuracy': 0.524,
'eval_precision': 0.4678362573099415, 'eval_recall': 0.7407407407407407,
'eval_f1': 0.5734767025089605, 'eval_runtime': 2.4012,
'eval_samples_per_second': 104.115, 'eval_steps_per_second': 0.833, 'epoch':
1.0}
Test metrics: {'eval_loss': 0.7307943105697632, 'eval_accuracy': 0.54,
'eval_precision': 0.5357142857142857, 'eval_recall': 0.7086614173228346,
'eval_f1': 0.6101694915254238, 'eval_runtime': 12.8892,
'eval_samples_per_second': 19.396, 'eval_steps_per_second': 0.155, 'epoch': 1.0}
```

```
[ ]: # save model checkpoint
timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
model_save_path = os.path.join(dir_models,
    f"{x_task}_{named_model}_{y_col}_{timestamp}")
trainer_obj.save_model(model_save_path)
print(f"Model checkpoint saved to: {model_save_path}")
# log experiment results
experiment_info = {
    "model_name": named_model,
    "learning_rate": learning_rate,
    "epochs": num_epochs,
```

```

        "batch_size": size_batch,
        "weight_decay": regularization_weight_decay,
        "x_task": x_task,
        "x_col": x_col,
        "y_col": y_col,
        "layers_to_unfreeze": layers_to_unfreeze}
model_info = gather_model_details(trained_model)
all_run_metrics = gather_all_run_metrics(
    trainer=trainer_obj,
    train_dataset=train_data_hf,
    val_dataset=val_data_hf,
    test_dataset=test_data_hf)
log_experiment_results_json(
    experiment_meta=experiment_info,
    model_details=model_info,
    run_metrics=all_run_metrics,
    log_file=log_filepath)
print(f"EXPERIMENT LOGGED TO: {log_filepath}")

```

Model checkpoint saved to:

/content/drive/MyDrive/266-final/models/multi\_answerdotai/ModernBERT-  
large\_binary\_complexity\_20250414\_015543

<IPython.core.display.HTML object>

EXPERIMENT LOGGED TO:

/content/drive/MyDrive/266-final/results/experiment\_runs.txt

```

[ ]: prediction_output = trainer_obj.predict(test_data_hf)
raw_predictions = prediction_output.predictions
true_labels = prediction_output.label_ids
preds = np.argmax(raw_predictions, axis=1)
mismatch_indices = np.where(preds != true_labels)[0]

error_rows = []
for idx in mismatch_indices:
    text_value = df_test.iloc[idx][x_col]
    true_label_val = true_labels[idx]
    pred_label_val = preds[idx]

    error_rows.append({
        "hf_index": idx,
        "text": text_value,
        "true_label": true_label_val,
        "predicted_label": pred_label_val
    })

error_df = pd.DataFrame(error_rows)

```

```

df_test_for_merge = df_test.copy()
df_test_for_merge["error_matching_prefix"] = df_test_for_merge[x_col].str[:50]
# df_test_for_merge.drop(columns=[x_col], inplace=True)

error_df["error_matching_prefix"] = error_df["text"].str[:50]
error_df = error_df.merge(
    df_test_for_merge,
    on="error_matching_prefix",
    how="left",
    suffixes=("", "_source"))

error_df.to_csv("bert-base-cased_mismatches.csv", index=False)

# print("Number of misclassified samples:", len(error_df))
print("\nMerged error_df with extra columns:")
# display(error_df.head(15))

# print("\nConfusion Matrix:")
# cm = confusion_matrix(true_labels, preds)
# plt.figure(figsize=(8, 6))
# sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=True,
#             xticklabels=["Predicted 0", "Predicted 1"],
#             yticklabels=["True 0", "True 1"])
# plt.xlabel('Predicted Label')
# plt.ylabel('True Label')
# plt.title('Confusion Matrix')
# plt.tight_layout()
# plt.show()
# print("confusion matrix metrics: \n", cm)
# error_save_path = os.path.join(dir_results,
#     ↳f"{x_task}_{named_model}_{y_col}_{x_col}_errors.csv")
# error_df.to_csv(error_save_path, index=False)
# print("Result saved to results directory.")
prediction_output = trainer_obj.predict(test_data_hf)
raw_predictions = prediction_output.predictions
true_labels = prediction_output.label_ids
preds = np.argmax(raw_predictions, axis=1)
df_test["true_label"] = true_labels
df_test["predicted_label"] = preds
df_test["is_incorrect"] = (df_test["predicted_label"] != df_test["true_label"])
df_test["avg_embedding"] = None
device = next(model.parameters()).device
for i in range(len(df_test)):
    text_value = df_test.iloc[i][x_col]
    e = tokenizer(text_value, return_tensors="pt", truncation=True,
        ↳max_length=512).to(device)
    with torch.no_grad():

```



```

        # emb = model.bert.embeddings.word_embeddings(e["input_ids"]).
        ↪mean(dim=1).squeeze().cpu().numpy()
        emb = model.model.embeddings.token_embeddings(e["input_ids"]).mean(dim=1).
        ↪squeeze().cpu().numpy()
        df_test.at[i, "avg_embedding"] = emb
cm = confusion_matrix(df_test["true_label"], df_test["predicted_label"])
plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=True,
        ↪xticklabels=["Predicted 0", "Predicted 1"], yticklabels=["True 0", "True 1"])
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.tight_layout()
plt.show()
# print("confusion matrix metrics:\n", cm)

df_plot = df_test.dropna(subset=["avg_embedding"]).copy()
embeddings = np.stack(df_plot["avg_embedding"].values)
pca = PCA(n_components=3)
reduced = pca.fit_transform(embeddings)
df_plot["pca_x"] = reduced[:,0]
df_plot["pca_y"] = reduced[:,1]
df_plot["pca_z"] = reduced[:,2]
colors = {
    ("bible", True): "red",
    ("bible", False): "orange",
    ("europarl", True): "yellow",
    ("europarl", False): "green",
    ("biomed", True): "blue",
    ("biomed", False): "purple"}
fig = plt.figure(figsize=(13,13))
ax = fig.add_subplot(111, projection='3d')
fig.set_facecolor("white")
ax.set_facecolor("white")
ax.xaxis._axinfo["grid"]["color"] = "gray"
ax.yaxis._axinfo["grid"]["color"] = "gray"
ax.zaxis._axinfo["grid"]["color"] = "gray"
for (corp, incorr), color in colors.items():
    subset = df_plot[(df_plot["corpus"]==corp) &
        ↪(df_plot["is_incorrect"]==incorr)]
    ax.scatter(subset["pca_x"], subset["pca_y"], subset["pca_z"], c=color,
        ↪s=20, alpha=0.8, label=f"{corp}, incorrect={incorr}")
ax.set_title("Average Embedding Value of Predictions, by Corpus", color="black")
ax.set_xlabel("PC1", color="black")
ax.set_ylabel("PC2", color="black")
ax.set_zlabel("PC3", color="black")
ax.legend(loc="best")

```

```

plt.show()
if "corpus" in df_test.columns:
    freqs = df_test["corpus"].value_counts()
    print("\nFrequency counts of corpus:", freqs)
    err_df = df_test[df_test["is_incorrect"]==True]
    corr_df = df_test[df_test["is_incorrect"]==False]
    err_counts = err_df["corpus"].value_counts()
    corr_counts = corr_df["corpus"].value_counts()
    print("\nCounts of corpus in misclassified:", err_counts)
    print("\nCounts of corpus in correctly classified:", corr_counts)
    print("\nProportions of corpus in misclassified:", err_counts/err_counts.
    ↪sum())
    print("\nProportions of corpus in correctly classified:", corr_counts/
    ↪corr_counts.sum())
    grouped_all = df_test.groupby("corpus")["avg_embedding"].apply(lambda x: np.
    ↪mean(np.stack(x.values), axis=0))
    grouped_err = err_df.groupby("corpus")["avg_embedding"].apply(lambda x: np.
    ↪mean(np.stack(x.values), axis=0))
    grouped_corr = corr_df.groupby("corpus")["avg_embedding"].apply(lambda x:
    ↪np.mean(np.stack(x.values), axis=0))
    # print("\nAvg embedding of each subcorpus overall:", grouped_all)
    # print("\nAvg embedding of each subcorpus misclassified:", grouped_err)
    # print("\nAvg embedding of each subcorpus correctly classified:",
    ↪grouped_corr)
    err_stack = np.stack(df_test[df_test["is_incorrect"]==True]["avg_embedding"].
    ↪values)
    corr_stack = np.stack(df_test[df_test["is_incorrect"]==False]["avg_embedding"].
    ↪values)
    # print("\nAvg embedding of records predicted incorrectly:", err_stack.
    ↪mean(axis=0))
    # print("Avg embedding of records predicted correctly:", corr_stack.
    ↪mean(axis=0))
    error_df = df_test[df_test["is_incorrect"]==True].copy()
    # error_df.to_csv("misclassified_with_all_columns.csv", index=False)
    # error_save_path = os.path.join(dir_results,
    ↪f"{x_task}_modernbert-base_{y_col}_{x_col}_errors.csv")
    df_test.to_csv(error_save_path, index=False)
    for corp in df_test["corpus"].unique():
        subset = df_test[df_test["corpus"]==corp]
        emb_true = subset[subset["is_incorrect"]==False]["avg_embedding"]
        emb_false = subset[subset["is_incorrect"]==True]["avg_embedding"]
        if len(emb_true)==0 or len(emb_false)==0:
            print(f"No valid data for subcorpus '{corp}'")
            continue
        p = np.mean(np.stack(emb_true.values), axis=0)
        q = np.mean(np.stack(emb_false.values), axis=0)

```

```

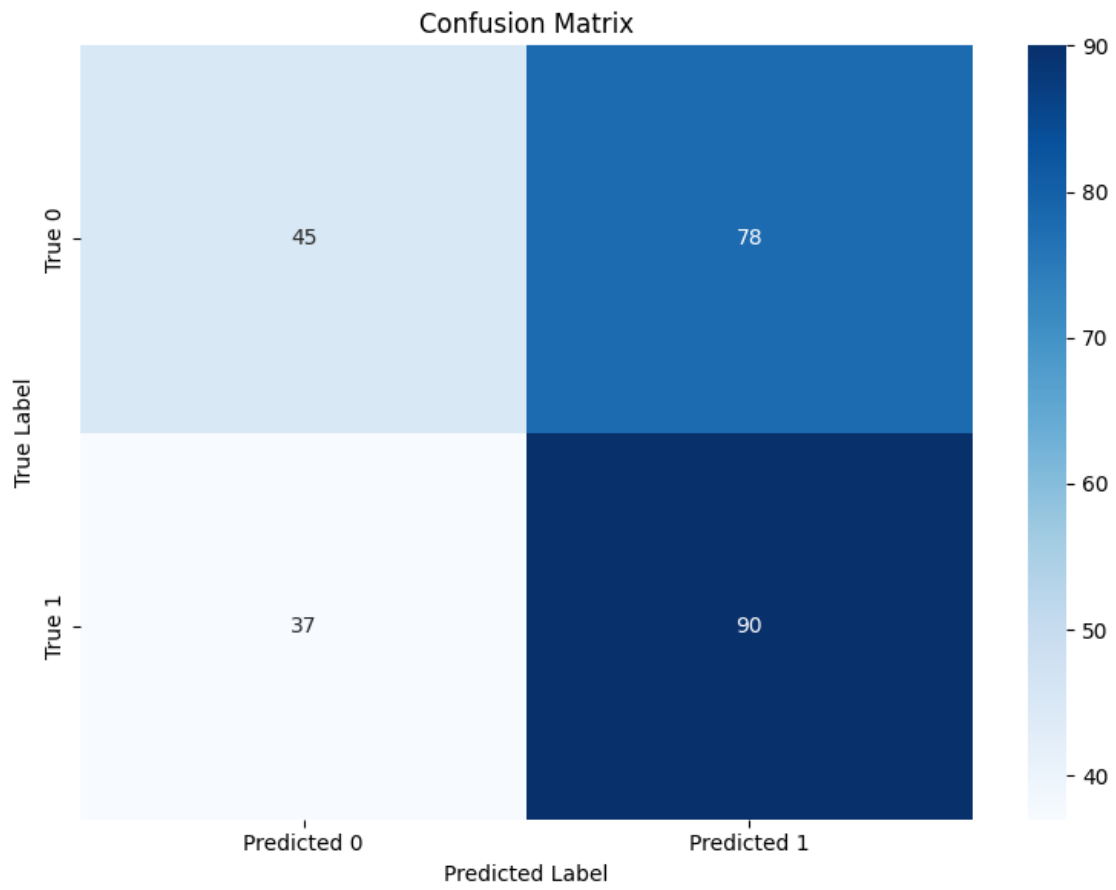
p_exp = np.exp(p - np.max(p))
q_exp = np.exp(q - np.max(q))
p_sum = p_exp.sum()
q_sum = q_exp.sum()
if p_sum<=0 or q_sum<=0:
    print(f"Cannot form valid distributions for subcorpus '{corp}'")
    continue
p_dist = p_exp / p_sum
q_dist = q_exp / q_sum
kl_pq = entropy(p_dist, q_dist)
kl_qp = entropy(q_dist, p_dist)
kl_sym = 0.5*(kl_pq + kl_qp)
print(f"Subcorpus '{corp}' symmetric KL divergence: {kl_sym}")
error_save_path = os.path.join(dir_results,
    ↪f"{x_task}_modernbert-large_{y_col}_{x_col}_errors.csv")
df_test.to_csv(error_save_path, index=False)
print("model results saved")


```

<IPython.core.display.HTML object>

Merged error\_df with extra columns:

<IPython.core.display.HTML object>





Frequency counts of corpus: corpus

bible 88

europarl 87

biomed 75

Name: count, dtype: int64

Counts of corpus in misclassified: corpus

europarl 47

bible 42

biomed 26

Name: count, dtype: int64

Counts of corpus in correctly classified: corpus

biomed 49

bible 46

europarl 40

Name: count, dtype: int64

Proportions of corpus in misclassified: corpus

europarl 0.408696

bible 0.365217

biomed 0.226087

Name: count, dtype: float64

Proportions of corpus in correctly classified: corpus

biomed 0.362963

bible 0.340741

europarl 0.296296

Name: count, dtype: float64

Subcorpus 'bible' symmetric KL divergence: 4.29344419874927e-06

Subcorpus 'biomed' symmetric KL divergence: 3.5779938949494955e-06

Subcorpus 'europarl' symmetric KL divergence: 4.058646832767972e-06

model results saved

**Result**

[ ]: